



University of Zurich^{UZH}

High Performance Computing Lecture 3

Douglas Potter

Cesare Cozza

Mark Eberlein

Where are
we?

We can run codes

- Submit on a SLURM queue

We can compile codes

- More on this today


Understanding Parallel

- Main topic today

GOAL: change a code

- Maybe write your own

Exercise:
Compute PI
in Parallel





Integration in Python

```
N=5
```

```
dx=1.0 / N
```

```
s = 0
```

```
for x in range(0, N):
```

```
    s += dx / ( 1 + ((x+0.5)*dx)**2 )
```

```
print(s*4)
```

```
dhcp-94-191:cpi$ python integrate.py
```

```
3.144925864003328
```



OpenMP Results (multi-threaded)

N=1,000,000,000

```
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=(nodes)
#SBATCH --ntasks-per-node=(processes)
#SBATCH --cpus-per-task=(threads)
#SBATCH --partition=normal
#SBATCH --constraint=mc

export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

srun cpi_openmp
```



OpenMP Results (multi-threaded)

N=1,000,000,000

```
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=36
#SBATCH --partition=normal
#SBATCH --constraint=mc

# 128 Eiger

export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

srun cpi_openmp
```



OpenMP Results (multi-threaded)

N=1,000,000,000

running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.1415926535899⁰¹ computed in 2.291 seconds
running on 3 threads: PI = 3.1415926535899⁶² computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589⁸²¹ computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589⁵⁹⁶ computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589⁶⁸² computed in 0.8988 seconds
running on 7 threads: PI = 3.141592653589⁶³ computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589⁷⁶⁹ computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589⁶⁵⁶ computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589⁷⁹⁴ computed in 0.5774 seconds
running on 11 threads: PI = 3.141592653589⁶⁶ computed in 0.5249 seconds
running on 12 threads: PI = 3.141592653589⁸⁶ computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589⁶⁵ computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589⁷⁸⁸ computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589⁸⁰⁵ computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589³² computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589³⁹ computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589¹⁴ computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589²⁶ computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589⁵⁵ computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589⁷⁷⁵ computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589⁸²³ computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589⁶⁶ computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589⁷⁹² computed in 0.2423 seconds
running on 25 threads: PI = 3.141592653589⁷⁸ computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589⁸³² computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589³⁵ computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589¹⁶ computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589¹⁹ computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589⁹³ computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589⁷⁴⁴ computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589⁷⁵⁸ computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589⁸⁰⁶ computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589⁹²⁶ computed in 0.1711 seconds
running on 35 threads: PI = 3.141592653589⁷⁹ computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589⁸²² computed in 0.1616 seconds



MPI Results

```
This is Process-1/36 running on nid00564
This is Process-3/36 running on nid00564
This is Process-5/36 running on nid00564
This is Process-6/36 running on nid00564
This is Process-8/36 running on nid00564
This is Process-11/36 running on nid00564
This is Process-13/36 running on nid00564
This is Process-14/36 running on nid00564
This is Process-17/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-22/36 running on nid00564
This is Process-23/36 running on nid00564
This is Process-25/36 running on nid00564
This is Process-26/36 running on nid00564
This is Process-28/36 running on nid00564
This is Process-29/36 running on nid00564
This is Process-31/36 running on nid00564
This is Process-33/36 running on nid00564
This is Process-35/36 running on nid00564
This is Process-0/36 running on nid00564
This is Process-2/36 running on nid00564
This is Process-4/36 running on nid00564
This is Process-7/36 running on nid00564
This is Process-9/36 running on nid00564
This is Process-10/36 running on nid00564
This is Process-12/36 running on nid00564
This is Process-15/36 running on nid00564
This is Process-16/36 running on nid00564
This is Process-19/36 running on nid00564
This is Process-20/36 running on nid00564
This is Process-21/36 running on nid00564
This is Process-24/36 running on nid00564
This is Process-27/36 running on nid00564
This is Process-30/36 running on nid00564
This is Process-32/36 running on nid00564
This is Process-34/36 running on nid00564
This program uses 36 processes
The number of intervals = 100000000
```

pi is approximately 3.1415926535898406
Error is 0.0000000000000475
wall clock time = 0.096959

OpenMP Result
running on 36 threads: PI = 3.141592653589822 computed in
0.1616 seconds

Two Nodes

(sometimes removed)

```
This is Process-62/72 running on nid01175
This is Process-22/72 running on nid01174
This is Process-63/72 running on nid01175
This is Process-23/72 running on nid01174
This is Process-64/72 running on nid01175
This is Process-24/72 running on nid01174
This is Process-65/72 running on nid01175
This is Process-25/72 running on nid01174
This is Process-66/72 running on nid01175
This is Process-26/72 running on nid01174
This is Process-67/72 running on nid01175
This is Process-29/72 running on nid01174
This is Process-68/72 running on nid01175
This is Process-30/72 running on nid01174
This is Process-69/72 running on nid01175
This is Process-31/72 running on nid01174
This is Process-70/72 running on nid01175
This is Process-32/72 running on nid01174
This is Process-71/72 running on nid01175
This is Process-33/72 running on nid01174
This is Process-39/72 running on nid01175
This is Process-34/72 running on nid01174
This is Process-40/72 running on nid01175
This is Process-35/72 running on nid01174
This is Process-41/72 running on nid01175
This is Process-0/72 running on nid01174
This is Process-43/72 running on nid01175
This is Process-45/72 running on nid01174
This is Process-46/72 running on nid01175
This is Process-11/72 running on nid01174
This is Process-13/72 running on nid01175
This is Process-47/72 running on nid01174
This is Process-19/72 running on nid01175
This is Process-51/72 running on nid01174
This is Process-20/72 running on nid01175
This is Process-59/72 running on nid01174
This is Process-27/72 running on nid01175
This is Process-28/72 running on nid01174
This program uses 72 processes
The number of intervals = 100000000
```

pi is approximately 3.1415926535898109,
Error is 0.0000000000000178
wall clock time = 0.024506

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=(nodes)
#SBATCH --ntasks-per-node=(processes)
#SBATCH --cpus-per-task=(threads)
#SBATCH --partition=normal
#SBATCH --constraint=mc

srun cpi_mpi
```




MPI Results

```
This is Process-1/36 running on nid00564
This is Process-3/36 running on nid00564
This is Process-5/36 running on nid00564
This is Process-6/36 running on nid00564
This is Process-8/36 running on nid00564
This is Process-11/36 running on nid00564
This is Process-13/36 running on nid00564
This is Process-14/36 running on nid00564
This is Process-17/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-22/36 running on nid00564
This is Process-23/36 running on nid00564
This is Process-25/36 running on nid00564
This is Process-26/36 running on nid00564
This is Process-28/36 running on nid00564
This is Process-29/36 running on nid00564
This is Process-31/36 running on nid00564
This is Process-33/36 running on nid00564
This is Process-35/36 running on nid00564
This is Process-0/36 running on nid00564
This is Process-2/36 running on nid00564
This is Process-4/36 running on nid00564
This is Process-7/36 running on nid00564
This is Process-9/36 running on nid00564
This is Process-10/36 running on nid00564
This is Process-12/36 running on nid00564
This is Process-15/36 running on nid00564
This is Process-16/36 running on nid00564
This is Process-19/36 running on nid00564
This is Process-20/36 running on nid00564
This is Process-21/36 running on nid00564
This is Process-24/36 running on nid00564
This is Process-27/36 running on nid00564
This is Process-30/36 running on nid00564
This is Process-32/36 running on nid00564
This is Process-34/36 running on nid00564
This program uses 36 processes
The number of intervals = 100000000
```

pi is approximately 3.1415926535898406
Error is 0.0000000000000475
wall clock time = 0.096959

OpenMP Result
running on 36 threads: PI = 3.141592653589822 computed in
0.1616 seconds

Two Nodes

(sometimes removed)

```
This is Process-62/72 running on nid01175
This is Process-22/72 running on nid01174
This is Process-63/72 running on nid01175
This is Process-23/72 running on nid01174
This is Process-64/72 running on nid01175
This is Process-24/72 running on nid01174
This is Process-65/72 running on nid01175
This is Process-25/72 running on nid01174
This is Process-66/72 running on nid01175
This is Process-26/72 running on nid01174
This is Process-67/72 running on nid01175
This is Process-29/72 running on nid01174
This is Process-68/72 running on nid01175
This is Process-30/72 running on nid01174
This is Process-69/72 running on nid01175
This is Process-31/72 running on nid01174
This is Process-70/72 running on nid01175
This is Process-32/72 running on nid01174
This is Process-71/72 running on nid01175
This is Process-33/72 running on nid01174
This is Process-39/72 running on nid01175
This is Process-34/72 running on nid01174
This is Process-40/72 running on nid01175
This is Process-35/72 running on nid01174
This is Process-41/72 running on nid01175
This is Process-42/72 running on nid01174
This is Process-43/72 running on nid01175
This is Process-44/72 running on nid01174
This is Process-45/72 running on nid01175
This is Process-46/72 running on nid01174
This is Process-47/72 running on nid01175
This is Process-11/72 running on nid01174
This is Process-48/72 running on nid01175
This is Process-13/72 running on nid01174
This is Process-49/72 running on nid01175
This is Process-20/72 running on nid01174
This is Process-50/72 running on nid01175
This is Process-27/72 running on nid01174
This is Process-59/72 running on nid01175
This is Process-28/72 running on nid01174
This program uses 72 processes
The number of intervals = 100000000
```

pi is approximately 3.1415926535898109,
Error is 0.0000000000000178
wall clock time = 0.024506

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=36 # 128 Eiger
#SBATCH --cpus-per-task=1
#SBATCH --partition=normal
#SBATCH --constraint=mc

srun cpi_mpi
```


Last
week

Same
results
in table
format

THREADS		SECONDS	
1	4.550	19	0.3053
2	2.291	20	0.2897
3	1.624	21	0.2768
4	1.258	22	0.2644
5	1.041	23	0.2528
6	0.8988	24	0.2423
7	0.7989	25	0.2326
8	0.7217	26	0.2237
9	0.6415	27	0.2154
10	0.5774	28	0.2077
11	0.5249	29	0.2006
12	0.4811	30	0.1939
13	0.4441	31	0.1876
14	0.4124	32	0.1818
15	0.3849	33	0.1763
16	0.3609	34	0.1711
17	0.3397	35	0.1662
18	0.3208	36	0.1616

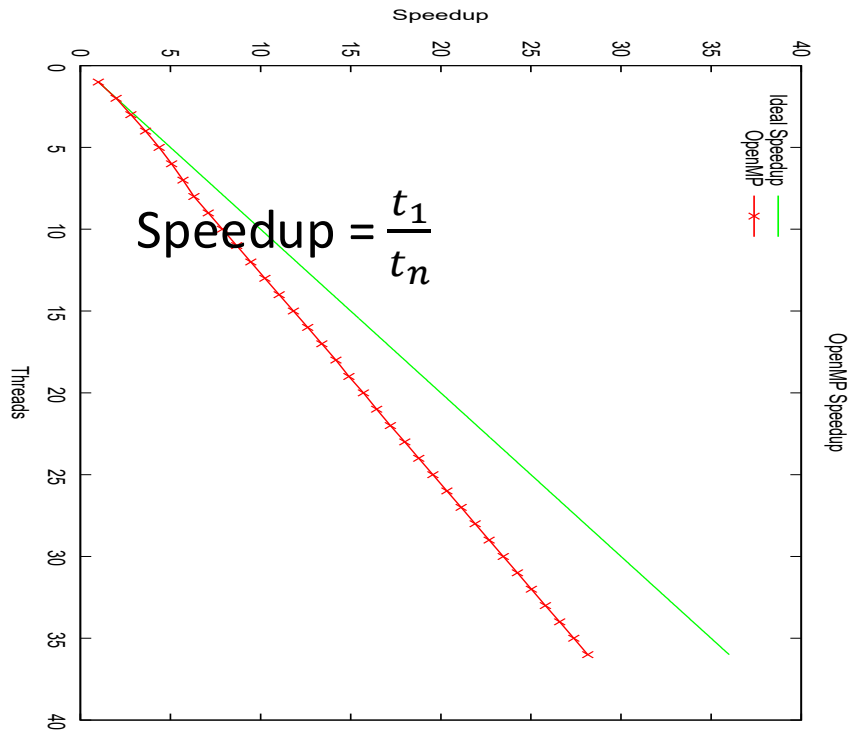
Timing Plot

```
set title 'OpenMP Timing'  
set xlabel 'Threads'  
set ylabel 'Time (seconds)'  
set key top right  
plot "cpi_openmp.dat"\  
u 1:2 w lp lw 2 t 'OpenMP'
```



“Speedup” Plot

```
set title 'OpenMP Speedup'  
set xlabel 'Threads'  
set ylabel 'Speedup'  
set key top left  
plot x lc 2 lw 2 t 'Ideal Speedup',\  
"cpi_openmp.dat" u 1:(4.55/$2)\  
w lp lc 1 lw 2 t 'OpenMP'
```



Parsing the Text (Results)



running on 1 threads: $\pi = 3.141592653589971$ computed in 4.55 seconds

1 4.550000	13 0.444100	25 0.232600
2 2.291000	14 0.412400	26 0.223700
3 1.624000	15 0.384900	27 0.215400
4 1.258000	16 0.360900	28 0.207700
5 1.041000	17 0.339700	29 0.200600
6 0.898800	18 0.320800	30 0.193900
7 0.798900	19 0.305300	31 0.187600
8 0.721700	20 0.289700	32 0.181800
9 0.641500	21 0.276800	33 0.176300
10 0.577400	22 0.264400	34 0.171100
11 0.524900	23 0.252800	35 0.166200
12 0.481100	24 0.242300	36 0.161600

```
./parse.pl <slurm-6188431.out
```

How to Plot the Results

```
dpotter@daint104:~/hpc1a> cat parse.pl
```

```
#!/usr/bin/perl
```

```
use strict;
```

```
while(<STDIN>) {
```

```
    if (/running on (\d+) threads: PI.*computed in ([0-9]*\.[0-9]+) seconds/) {
```

```
        printf("%2d %f\n", $1, $2);
```

```
    }
```

```
}
```

```
dpotter@daint104:~/hpc1a> cat speedup.gp
```

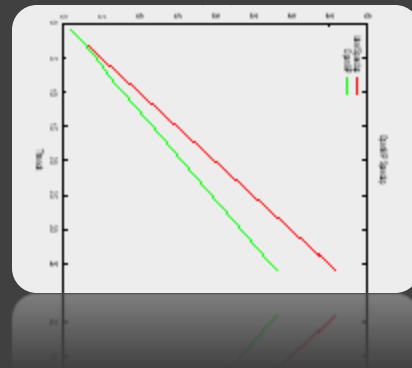
```
set title 'OpenMP Speedup'
```

```
set xlabel 'Threads'
```

```
set ylabel 'Speedup'
```

```
set key top left
```

```
plot x t 'Ideal Speedup', "cpi_openmp.dat" u 1:(4.55/$2) w l t 'OpenMP'
```



```
./parse.pl <slurm-6188431.out >cpi_openmp.dat
```

```
gnuplot> set output "speedup.eps"
```

```
gnuplot> set terminal postscript enhanced solid color
```

```
gnuplot> load "speedup.gp"
```

```
gnuplot> set terminal x11
```

```
gnuplot> set output
```

Parsing with Python


```
dpotter@daint104:~/hpc1a> cat parse.py
#!/usr/bin/env python
import re, sys

for line in re.finditer(
    r"running on ([0-9]+) threads: PI.*computed in ([0-9]*\.[0-9]+) seconds",
    sys.stdin.read(), re.MULTILINE):
    print("{} {}".format(line.group(1), line.group(2)))
```

```
dpotter@daint104:~/hpc1a> cat speedup.gp
set title 'OpenMP Speedup'
set xlabel 'Threads'
set ylabel 'Speedup'
set key top left
plot x t 'Ideal Speedup', "cpi_openmp.dat" u 1:(4.55/$2) w l t 'OpenMP'
```

```
./parse.py <slurm-6188431.out >cpi_openmp.dat
```

```
gnuplot> set output "speedup.eps"
gnuplot> set terminal postscript enhanced solid color
gnuplot> load "speedup.gp"
gnuplot> set terminal x11
gnuplot> set output
```





Parsing Process

```
[dpotter@ela4 hpc1a]$ head -n 36 slurm-6188431.out
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
[dpotter@ela4 hpc1a]$ python parse.py <slurm-6188431.out
1 4.55
2 2.291
3 1.624
4 1.258
5 1.041
6 0.8988
7 0.7989
8 0.7217
9 0.6415
10 0.5774
11 0.5249
12 0.4811
13 0.4441
14 0.4124
15 0.3849
16 0.3609
17 0.3397
18 0.3208
19 0.3053
20 0.2897
21 0.2768
22 0.2644
23 0.2528
24 0.2423
25 0.2326
26 0.2237
27 0.2154
28 0.2077
29 0.2006
30 0.1939
31 0.1876
32 0.1818
33 0.1763
34 0.1711
35 0.1662
36 0.1616
```


More on Compilers

High Performance Computing





Compiler Suites

[Cray Compiler](#) ([PrgEnv-cray](#))

[GNU Compiler](#) ([PrgEnv-gnu](#))

- Free for everyone

[Intel Compiler](#) ([PrgEnv-intel](#))

- [Free for Students](#)

[Portland Compiler](#) ([PrgEnv-pgi](#))

- Important for OpenACC (GPU)

[clang](#) (Apple and Cray)

- Open Source & GNU Compatible

[Visual Studio](#) (Windows)

- [Free for us through Microsoft Imagine](#)

```

#include <stdio.h>
#include <sys/time.h>

static long steps = 1000000000;

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;
    char *p;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}

```

cpi.c
(serial
version)



man gettimeofday

GETTIMEOFDAY(2)

System Calls Manual

GETTIMEOFDAY(2)

NAME

gettimeofday, **settimeofday** – get/set date and time

SYNOPSIS

#include <sys/time.h>

int

gettimeofday(struct timeval *restrict tp, void *restrict tzp);

int

settimeofday(const struct timeval *tp, const struct timezone *tzp);

DESCRIPTION

The system's notion of the current Greenwich time and the current time zone is obtained with the **gettimeofday()** call, and set with the **settimeofday()** call. The time is expressed in seconds and microseconds since midnight (0 hour), January 1, 1970. The resolution of the system clock is hardware dependent, and the time may be updated continuously or in “ticks.”

```

#include <stdio.h>
#include <sys/time.h>

static long steps = 1000000000;

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;
    char *p;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

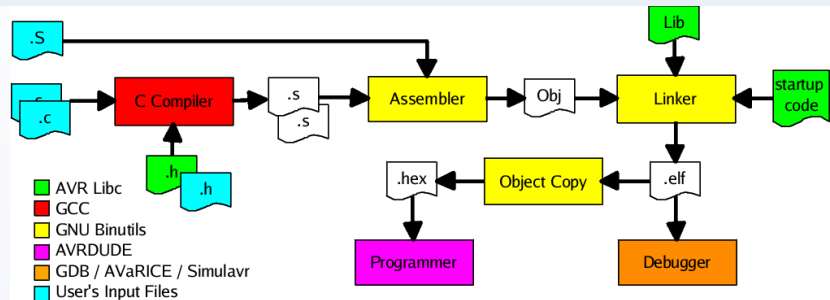
    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}

```

cpi.c
(serial
version)

Output Name

```
dhcp-94-191:cpu$ ls
cpu.c
dhcp-94-191:cpu$ gcc cpu.c
dhcp-94-191:cpu$ ls
a.out cpu.c
dhcp-94-191:cpu$ ./a.out
PI = 3.141592653589971 computed in 11.9 seconds
dhcp-94-191:cpu$ rm a.out
dhcp-94-191:cpu$ ls
cpu.c
dhcp-94-191:cpu$ gcc -o cpu cpu.c
dhcp-94-191:cpu$ ls
cpu cpu.c
dhcp-94-191:cpu$ ./cpu
PI = 3.141592653589971 computed in 11.82 seconds
```



Compiler Optimization

Semicolon (command separator)
Multiple command on one line

```
dhcp-94-191:cpu$ gcc -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 11.84 seconds
dhcp-94-191:cpu$ gcc -O -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 2.505 seconds
dhcp-94-191:cpu$ gcc -O0 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 11.82 seconds
dhcp-94-191:cpu$ gcc -O1 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 2.497 seconds
dhcp-94-191:cpu$ gcc -O2 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 1.19 seconds
dhcp-94-191:cpu$ gcc -O3 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 1.174 seconds
dhcp-94-191:cpu$ gcc -O3 -ffast-math -o cpi cpi.c ; ./cpi
PI = 3.141592653589652 computed in 0.6 seconds
dhcp-94-191:cpu$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c ; ./cpi
PI = 3.141592653589729 computed in 0.5518 seconds
3.14159265358979323846 (real value of pi)
```

> 20 Times
Improvement

What do the levels mean?

-O1 Optimize.

- Optimizing compilation takes somewhat more time, and a lot more memory for a large function.

-O2 Optimize even more.

- GCC performs nearly all supported optimizations that do not involve a space-speed tradeoff.
- The compiler does not perform loop unrolling or function inlining when you specify -O2.
- As compared to -O, this option **increases both compilation time** and the performance of the generated code.

-O3 Optimize yet more.

- -O3 turns on all optimizations specified by -O2 and also turns on the -finline-functions, -funswitch-loops, -fpredictive-commoning, -fgcse-after-reload and -ftree-vectorize options.

-O0 No Optimization.

- Reduce compilation time and make debugging produce the expected results.
- **This is the default.**

Fast Math

Associative Property of Addition

$$(a + b) + c = a + (b + c)$$

Distributive Property of Multiplication

$$a(b + c) = ab + ac$$

Sometimes reordering can
improve performance

Fast Math (the Truth)

Associative Property of Addition

$$(a + b) + c \neq a + (b + c)$$

Distributive Property of Multiplication

$$a(b + c) \neq ab + ac$$

Sometimes reordering can
change the result

Beware of Associative Math

```
dhcp-94-191:cpi$ python
```

```
>>> a=12
```

```
>>> b=12
```

```
>>> a==b
```

```
True
```

```
>>> a=0.1 + 0.2 + 0.3
```

```
>>> b=0.1 +(0.2 + 0.3)
```

```
>>> a==b
```

```
False
```

```
>>> print(a)
```

```
0.6
```

```
>>> print(b)
```

```
0.6
```

```
>>> print(f"{a:.17f}")
```

```
0.60000000000000009
```

```
>>> print(f"{b:.17f}")
```

```
0.59999999999999998
```

```
>>> c=0.6
```

```
>>> print(f"{c:.17f}")
```

```
0.59999999999999998
```

“Modules” in C

```
#include <stdio.h>
#include <sys/time.h>

static long steps = 1000000000;

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}

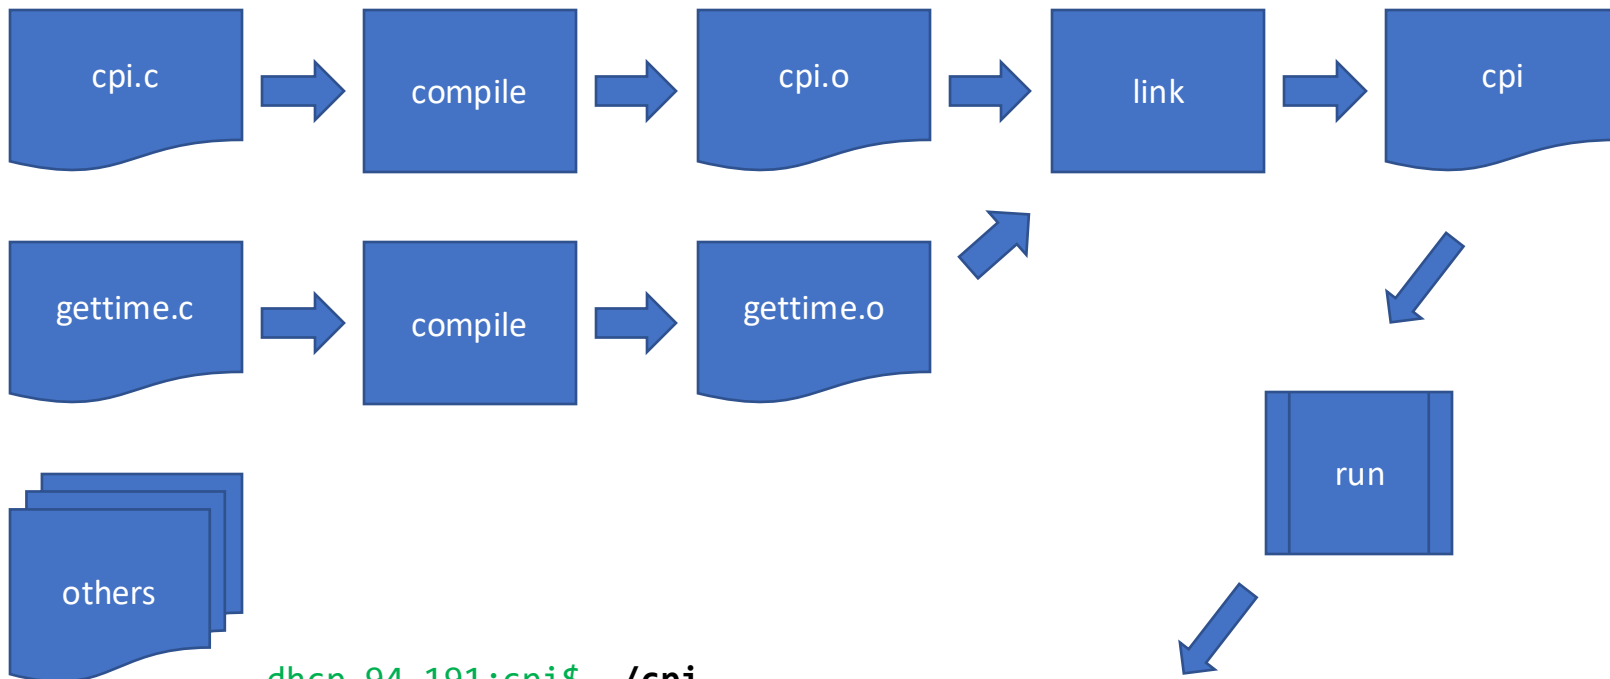
int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Separate Compilation



```
dhcp-94-191:cpi$ ./cpi
```

```
PI = 3.141592653589971 computed in 1.174 seconds
```



gettime.c

```
#include <sys/time.h>

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}
```



cpi.c

```
#include <stdio.h>

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Compiling with two files

```
dhcp-94-191:cpi2$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c gettime.c
```

```
cpi.c: In function 'main':
```

```
cpi.c:13:20: warning: implicit declaration of function 'getTime'; did you mean 'getline'? [-Wimplicit-function-declaration]
```

```
    double start = getTime();
```

```
                ^~~~~~
```

```
                getline
```




University of
Zurich ^{UZH}

gettime.h

```
double getTime(void);
```



cpi.c

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();
    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

gettime.h

```
double getTime(void);
```



gettime.c

```
#include <sys/time.h>
#include "gettime.h"

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}
```

```
gettime.h
```

```
double getTime(void);
```



Compile and Link Separately

```
dhcp-94-191:pi2$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c gettimeofday.c
dhcp-94-191:pi2$ ls
cpi          cpi.c      gettimeofday.c gettimeofday.h
dhcp-94-191:pi2$ rm cpi
```

```
dhcp-94-191:pi2$ ls
cpi.c      gettimeofday.c gettimeofday.h
dhcp-94-191:pi2$ gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
dhcp-94-191:pi2$ ls
cpi.c      cpi.o      gettimeofday.c gettimeofday.h gettimeofday.o
dhcp-94-191:pi2$ gcc -O3 -ffast-math -mavx2 -c -o gettimeofday.o gettimeofday.c
dhcp-94-191:pi2$ ls
cpi.c      cpi.o      gettimeofday.c gettimeofday.h gettimeofday.o
dhcp-94-191:pi2$ gcc -o cpi cpi.o gettimeofday.o
dhcp-94-191:pi2$ ./cpi
PI = 3.141592653589729 computed in 0.567 seconds
```



Makefile

```

cpu
    : cpu.o gettime.o
    gcc -o cpu cpu.o gettime.o

cpu.o
    : cpu.c gettime.h
    gcc -O3 -ffast-math -mavx2 -c -o cpu.o cpu.c

gettime.o
    : gettime.c gettime.h
    gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c

clean:
    rm -f cpu cpu.o gettime.o
```



Compile and Link Separately

```
dhcp-94-191:cp12$ make clean
rm -f cpi cpi.o gettime.o
dhcp-94-191:cp12$ make
gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c
gcc -o cpi cpi.o gettime.o
dhcp-94-191:cp12$ make
make: `cpi' is up to date.
dhcp-94-191:cp12$ touch cpi.c
dhcp-94-191:cp12$ make
gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
gcc -o cpi cpi.o gettime.o
```

“touch” – change modification time.
Same as if you edited the file.



Makefile

```

api      : api.o gettime.o
            gcc -o api api.o gettime.o

api.o    : api.c gettime.h
            gcc -O3 -ffast-math -mavx2 -c -o api.o api.c

gettime.o: gettime.c gettime.h
            gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c

clean:
            rm -f api api.o gettime.o
```



Makefile with default rules

```
cp            : cpi.o gettime.o
```

```
cpi.o        : cpi.c gettime.h
```

```
gettime.o:   gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpi cpi.o gettime.o
```

```
dhcp-94-191:cp12$ make clean
```

```
rm -f cpi cpi.o gettime.o
```

```
dhcp-94-191:cp12$ make
```

```
cc -c -o cpi.o cpi.c
```

```
cc -c -o gettime.o gettime.c
```

```
cc cpi.o gettime.o -o cpi
```




Makefile with default rules

```
CFLAGS=-O3 -ffast-math -mavx2
```

```
CC=gcc
```

```
cpu      : cpu.o gettime.o
```

```
cpu.o    : cpu.c gettime.h
```

```
gettime.o: gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpu cpu.o gettime.o
```

```
dhcp-94-191:cpu2$ make clean
```

```
rm -f cpu cpu.o gettime.o
```

```
dhcp-94-191:cpu2$ make
```

```
gcc -O3 -ffast-math -mavx2 -c -o cpu.o cpu.c
```

```
gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c
```

```
gcc  cpu.o gettime.o -o cpu
```



Message Passing

Why use MPI?

One Node
isn't enough

- Maybe you have run out of memory
- Maybe it takes too long to calculate

Alternatives
do exist

- Partitioned Global Address Space (PGAS)
- HPX (High Performance ParallelX)
- Charm++

MPI is still
omnipresent

- It has been around since 1991
- Version 3.1 was approved in 2015
- Version 4.0 release candidate now (2021)

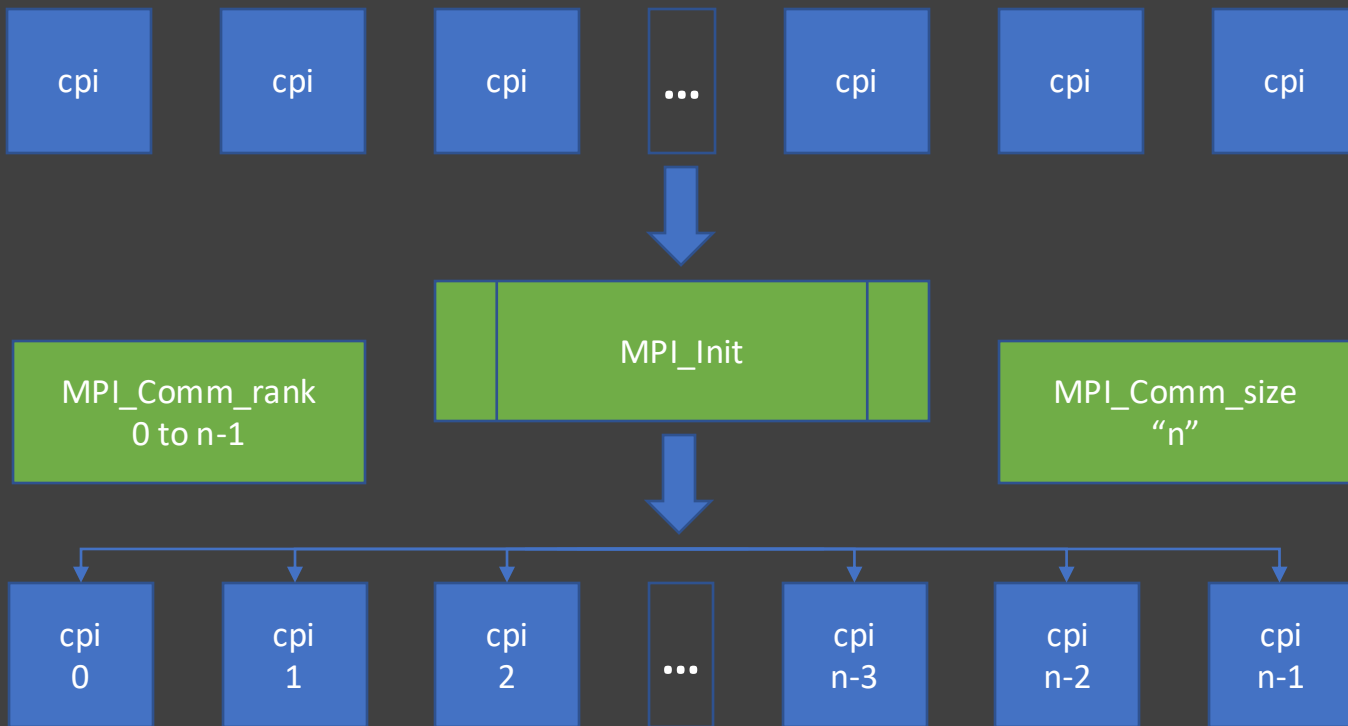


MPI Process Layout

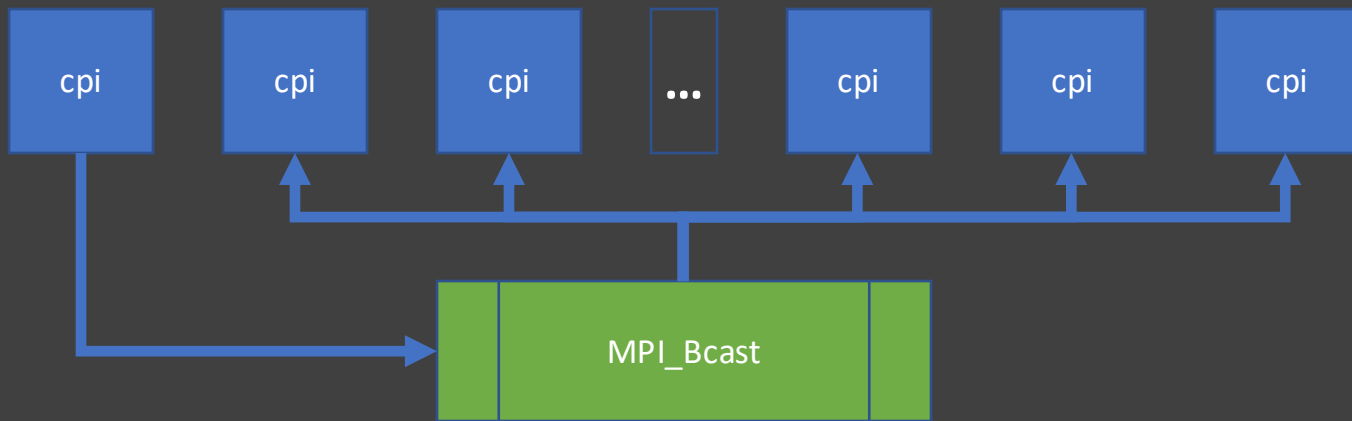




MPI Initialization



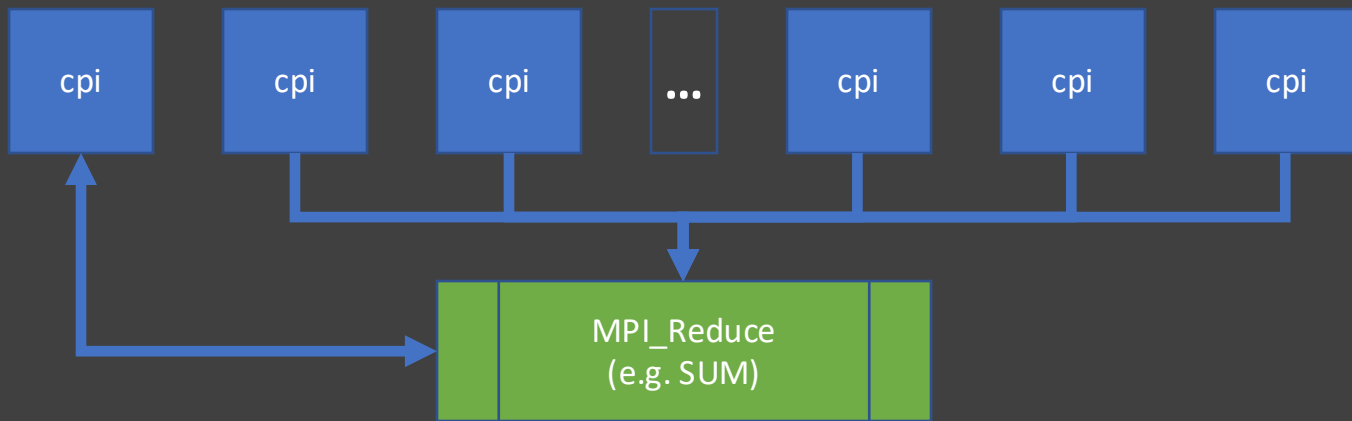
MPI Broadcast



All MPI Ranks Call MPI_Bcast together.
Everyone gets the value from (usually) rank 0.

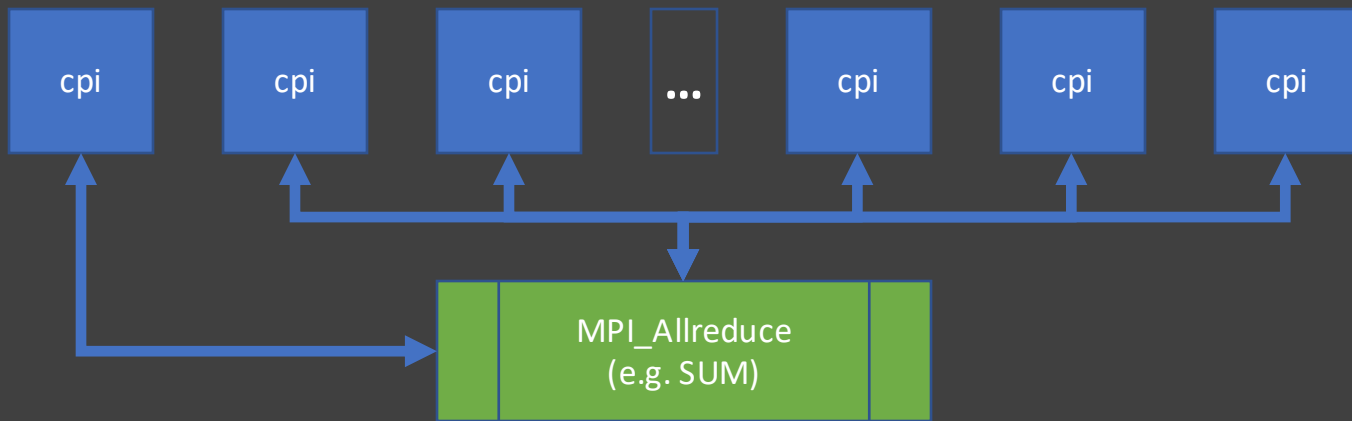


MPI Reduce



All MPI Ranks Call MPI_Reduce together.
All values from all ranks are summed together.
A single rank (usually rank 0) gets the result.



MPI All Reduce



All MPI Ranks Call MPI_Reduce together.
All values from all ranks are summed together.
All ranks get the result.

Partial Reduction (e.g., sum)

Rank	0	1	2	3	4	5	6	7
Value	4	3	8	7	2	6	3	1
Scan	4	7	15	22	24	30	33	34
Exscan	0	4	7	15	22	24	30	33





The MPI version of cpi

```
MPI_Init(&argc,&argv);           /* Connect processes to each other */
MPI_Comm_size(MPI_COMM_WORLD,&numprocs); /* Get total number of processes */
MPI_Comm_rank(MPI_COMM_WORLD,&myid);    /* Rank of this process */
MPI_Get_processor_name(name, &resultlen);
```

```
printf("This is Process-%d/%d running on %s \n",myid,numprocs,name);
MPI_Barrier(MPI_COMM_WORLD);
```

```
if(myid == 0) {
    printf("This program uses %d processes\n", numprocs);
    n = 1000000000;
}
```

```
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); /* Send n from 0 to all others */
```

```
sum = 0.0;
h = 1.0/n;
for (i=myid+0.5; i<n; i+=numprocs) {
    sum += dx_arctan(i*h);
}
mypi = 4.0*h*sum;
```

```
/* Compute the Derivative of ArcTan */
double dx_arctan(double x) {
    return 1.0 / (1.0 + x*x);
}
```

```
/* Add all partial sums to each other and send to rank 0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

```
if (myid == 0) {
    printf("pi is approximately %.16f\n",pi);
}
```



```
MPI_Finalize(); /* Disconnect all processes */
```





MPI File write ordered en	MPI Put	MPI Publish name
MPI Comm c2f	MPI File write shared	MPI Put
MPI Comm call errhandler	MPI Finalize	MPI Query thread
MPI Comm compare	MPI Finalized	MPI Raccumulate
MPI Comm connect	MPI Free mem	MPI Recv
MPI Comm create	MPI Gather	MPI Recv init
MPI Comm create errhandler	MPI Gatherv	MPI Reduce
MPI Comm create group	MPI Get	MPI Reduce local
MPI Comm create keyval	MPI Get accumulate	MPI Reduce scatter
MPI Comm delete attr	MPI Get address	MPI Reduce scatter block
MPI Comm disconnect	MPI Get count	MPI Register datarep
MPI Comm dup	MPI Get elements	MPI Request c2f
MPI Comm dup with info	MPI Get elements x	MPI Request f2c
MPI Comm f2c	MPI Get library version	MPI Request free
MPI Comm free	MPI Get processor name	MPI Request get status
MPI Comm free keyval	MPI Get version	MPI Rget
MPI Comm get attr	MPI Graph create	MPI Rget accumulate
MPI Comm get errhandler	MPI Graph get	MPI Rput
MPI Comm get info	MPI Graph map	MPI Rsend
MPI Comm get name	MPI Graph neighbors	MPI Rsend init
MPI Comm get name get	MPI Graph neighbors count	MPI Scan
MPI Comm group	MPI Graphdims get	MPI Scatter
MPI Comm idup	MPI Grequest complete	MPI Scatterv
MPI Comm join	MPI Grequest start	MPI Send
MPI Comm rank	MPI Group c2f	MPI Send init
MPI Comm remote group	MPI Group compare	MPI Sendrecv
MPI Comm remote size	MPI Group difference	MPI Sendrecv replace
MPI Comm set attr	MPI Group excl	MPI Sizeof
MPI Comm set errhandler	MPI Group f2c	MPI Ssend
MPI Comm set info	MPI Group free	MPI Ssend init
MPI Comm set name	MPI Group incl	MPI Start
MPI Comm size	MPI Group intersection	MPI Startall
MPI Comm spawn	MPI Group range excl	MPI Status c2f
MPI Comm spawn multiple	MPI Group range incl	MPI Status f2c
MPI Comm split	MPI Group rank	MPI Status set cancelled
MPI Comm split type	MPI Group size	MPI Status set elements
MPI Comm test inter	MPI Group translate ranks	MPI Status set elements x
MPI Compare and swap	MPI Group union	MPI T category changed
MPI Dims create	MPI Iallgather	MPI T category get categories
MPI Dist graph create	MPI Iallgatherv	MPI T category get cvars
MPI Dist graph create adjacent	MPI Iallreduce	MPI T category get info
MPI Dist graph neighbors	MPI Ialltoall	MPI T category get num
MPI Dist graph neighbors count	MPI Ialltoallv	MPI T category get pvars

MPI File call errhandler	MPI Ineighbor allgather	MPI T init thread
MPI File close	MPI Ineighbor allgatherv	MPI T pvar_get info
MPI File create errhandler	MPI Ineighbor alltoall	MPI T pvar_get num
MPI File delete	MPI Ineighbor alltoallv	MPI T pvar_handle alloc
MPI File f2c	MPI Ineighbor alltoallw	MPI T pvar_handle free
MPI File get amode	MPI Info c2f	MPI T pvar_read
MPI File get atomicity	MPI Info create	MPI T pvar_readreset
MPI File get byte offset	MPI Info delete	MPI T pvar_reset
MPI File get errhandler	MPI Info dup	MPI T pvar_session create
MPI File get group	MPI Info env	MPI T pvar_session free
MPI File get info	MPI Info f2c	MPI T pvar_start
MPI File get position	MPI Info free	MPI T pvar_stop
MPI File get position shared	MPI Info get	MPI T pvar_write
MPI File get size	MPI Info get nkeys	MPI Test
MPI File get type extent	MPI Info get nthkey	MPI Test cancelled
MPI File get view	MPI Info get valuelen	MPI Testall
MPI File iread	MPI Info set	MPI Testany
MPI File iread all	MPI Init	MPI Testsome
MPI File iread at	MPI Init thread	MPI Topo_test
MPI File iread at all	MPI Initialized	MPI Type c2f
MPI File iread shared	MPI Intercomm create	MPI Type commit
MPI File iwrite	MPI Intercomm merge	MPI Type contiguous
MPI File iwrite all	MPI Iprobe	MPI Type create darray
MPI File iwrite at	MPI Irecv	MPI Type create f90 co
MPI File iwrite at all	MPI Ireduce	MPI Type create f90 int
MPI File iwrite shared	MPI Ireduce scatter	MPI Type create f90 res
MPI File open	MPI Ireduce scatter block	MPI Type create hindexe
MPI File preallocate	MPI Irsend	MPI Type create hindexe
MPI File read	MPI Is_thread_main	MPI Type create hvecto
MPI File read all	MPI Iscan	MPI Type create indexe
MPI File read all begin	MPI Iscatter	MPI Type create keyval
MPI File read all end	MPI Iscatterv	MPI Type create resized
MPI File read at	MPI Isend	MPI Type create struct
MPI File read at all	MPI Issend	MPI Type create subarra
MPI File read at all begin	MPI Keyval create	MPI Type delete attr
MPI File read at all end	MPI Keyval free	MPI Type dup
MPI File read ordered	MPI Lookup_name	MPI Type extent
MPI File read ordered begin	MPI Message c2f	MPI Type f2c
MPI File read ordered end	MPI Message f2c	MPI Type free
MPI File read ordered	MPI Mprobe	MPI Type free keyval
MPI File read ordered	MPI Mrecv	MPI Type get attr
MPI File read ordered	MPI Neighbor allgather	MPI Type get contents
MPI File read ordered	MPI Neighbor allgatherv	MPI Type get envelope
MPI File read ordered	MPI Neighbor alltoall	MPI Type get extent
MPI File read ordered	MPI Neighbor alltoallv	MPI Type get extent x
MPI File read ordered	MPI Neighbor alltoallw	MPI Type get name
MPI File read ordered	MPI Op c2f	MPI Type get true exte
MPI File read ordered	MPI Op commutative	MPI Type get true exte
MPI File read ordered	MPI Op create	MPI Type indexed
MPI File read ordered	MPI Op f2c	MPI Type hvector
MPI File read ordered	MPI Op free	MPI Type indexed
MPI File read ordered	MPI Open_port	MPI Type lb
MPI File read ordered	MPI Pack	MPI Type match size
MPI File read ordered	MPI Pack external	MPI Type set attr



OpenMP



High Performance Computing

The Serial Version of cpi

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

Start of OpenMP Version

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    #pragma omp parallel for
    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Need to Update our Makefile

```
CFLAGS=-Wall -O3 -ffast-math -mavx2 -fopenmp
```

```
CC=gcc
```

```
cpu : cpu.o gettime.o
```

```
cpu.o : cpu.c gettime.h
```

```
gettime.o : gettime.c gettime.h
```

```
clean:
```

```
rm -f cpu cpu.o gettime.o
```



Not quite good enough

```
dhcp-94-191:cp13$ make
gcc -Wall -O3 -ffast-math -mavx2 -fopenmp -c -o cpi.o cpi.c
gcc -Wall -O3 -ffast-math -mavx2 -fopenmp -c -o gettime.o gettime.c
gcc cpi.o gettime.o -o cpi
Undefined symbols for architecture x86_64:
  "_GOMP_parallel", referenced from:
      _main in cpi.o
  "_omp_get_num_threads", referenced from:
      _main.omp_fn.0 in cpi.o
  "_omp_get_thread_num", referenced from:
      _main.omp_fn.0 in cpi.o
ld: symbol(s) not found for architecture x86_64
collect2: error: ld returned 1 exit status
make: *** [cpi] Error 1
```




Makefile need linker flags too

```
CFLAGS=-Wall -O3 -ffast-math -mavx2 -fopenmp
```

```
LDFLAGS=-fopenmp
```

```
CC=gcc
```

```
cpu                : cpu.o gettime.o
```

```
cpu.o              : cpu.c gettime.h
```

```
gettime.o          : gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpu cpu.o gettime.o
```

Better. Not good, but better.

```
dhcp-94-191:cp13$ make  
gcc -fopenmp cpi.o gettimeofday.o -o cpi  
dhcp-94-191:cp13$ ./cpi  
PI = 0.5675882184166633 computed in 0.2785 seconds
```

(Laptop)

This is pi

This is not pi

```
dhcp-94-191:cp13$ OMP_NUM_THREADS=1 ./cpi  
PI = 3.141592653589729 computed in 0.5467 seconds  
dhcp-94-191:cp13$ OMP_NUM_THREADS=2 ./cpi  
PI = 1.287002217586625 computed in 0.2806 seconds  
dhcp-94-191:cp13$ OMP_NUM_THREADS=4 ./cpi  
PI = 0.9799146525073925 computed in 0.2782 seconds  
dhcp-94-191:cp13$ OMP_NUM_THREADS=4 ./cpi  
PI = 0.5675882184166633 computed in 0.279 seconds  
dhcp-94-191:cp13$ OMP_NUM_THREADS=4 ./cpi  
PI = 0.9799146525073925 computed in 0.2772 seconds
```

Could be worse though (ARM)

```
$ ./cpi # serial  
PI = 3.141592653589839 computed in 1.919 seconds  
$ ./cpi # OpenMP  
PI = 0.1722031954746213 computed in 7.018 seconds  
$ OMP_NUM_THREADS=1 ./cpi  
PI = 3.141592653589839 computed in 6.849 seconds  
$ OMP_NUM_THREADS=2 ./cpi  
PI = 1.545778454564867 computed in 4.479 seconds  
$ OMP_NUM_THREADS=4 ./cpi  
PI = 0.7602211683355252 computed in 2.424 seconds  
$ OMP_NUM_THREADS=8 ./cpi  
PI = 0.4593441363629033 computed in 4.056 seconds  
$ OMP_NUM_THREADS=16 ./cpi  
PI = 0.193699904454339 computed in 5.435 seconds
```

Where is the problem?

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    #pragma omp parallel for
    for (i=0; i < steps; i++) {
        double x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

Where is the problem?

```
#include <stdio.h>
#include "gettime.h"
```

```
static long steps = 1000000000;
```

```
int main (int argc, const char *argv[]) {
    int i;
    double pi;
```

```
    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();
```

Like
MPI_Reduce

```
#pragma omp parallel for reduction(+ : sum)
```

```
    for (i=0; i < steps; i++) {
        double x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
```

Okay if declared inside loop

All others must be "handled"

```
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

Loop variable is fine

With “reduction” it works!

```
$ ./cpi # serial
```

```
PI = 3.141592653589839 computed in 1.919 seconds
```

```
$ ./cpi # OpenMP
```

```
PI = 3.141592653589825 computed in 0.1358 seconds
```

```
$ OMP_NUM_THREADS=1 ./cpi
```

```
PI = 3.141592653589839 computed in 1.913 seconds
```

```
$ OMP_NUM_THREADS=2 ./cpi
```

```
PI = 3.141592653589855 computed in 0.9901 seconds
```

```
$ OMP_NUM_THREADS=4 ./cpi
```



```
PI = 3.141592653589803 computed in 0.5135 seconds
```

```
$ OMP_NUM_THREADS=8 ./cpi
```

```
PI = 3.141592653589804 computed in 0.2629 seconds
```

```
$ OMP_NUM_THREADS=16 ./cpi
```

```
PI = 3.141592653589738 computed in 0.1437 seconds
```




Parallel Pitfalls

What you're glad that
you don't need to know
(too much)




Critical Section: $\text{sum} += \text{value}$





Solution: Locking




Deadlock

- Think for some time
- Pick up left fork
- Pick up right fork
- Eat for some time
- Put down right fork
- Put down left fork
- Continue Thinking



Edsger W. Dijkstra



Requirements for a Deadlock

Mutual Exclusion

- Resource cannot be shared

Hold and Wait

- Must hold a resource while waiting for another

No Pre-emption

- Held resources cannot be given away


Circular Wait

- The wait dependency must be circular

Starvation

(No Hold and Wait)



- Think for some time
- Pick up left fork
- Right fork not available
- ... or left fork not available
- Put down left fork
- Think for a while
- Eventually try again



Resource Hierarchy (No Circular Wait)

- Number each resource
 - Confucius is different
- Choose resources in order
- Cannot deadlock
- Cannot starve?







Load Balancing

How to divide
up the work


Divide 100 rectangles
into four domains




Option 1




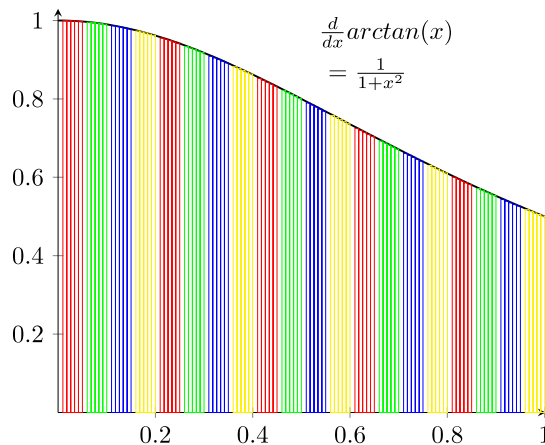
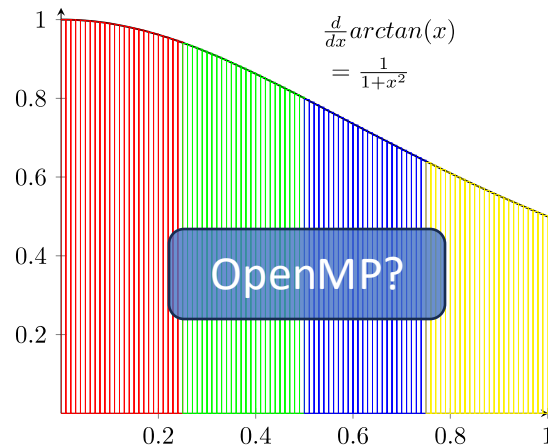

Option 2



Option 3



Which is the best choice?





The MPI version of cpi

```
MPI_Init(&argc,&argv); /* Connect processes to each other */
MPI_Comm_size(MPI_COMM_WORLD,&numprocs); /* Get total number of processes */
MPI_Comm_rank(MPI_COMM_WORLD,&myid); /* Rank of this process */
MPI_Get_processor_name(name, &resultlen);

printf("This is Process-%d/%d running on %s \n",myid,numprocs,name);
MPI_Barrier(MPI_COMM_WORLD);


if(myid == 0) {
    printf("This program uses %d processes\n", numprocs);
    n = 1000000000;
}

MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); /* Send n from 0 to all others */

sum = 0.0;
h = 1.0/n;
for (i=myid+0.5; i<n; i+=numprocs) {
    sum += dx_arctan(i*h);
}
mypi = 4.0*h*sum;

/* Add all partial sums to each other and send to rank 0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if (myid == 0) {
    printf("pi is approximately %.16f\n",pi);
}
MPI_Finalize(); /* Disconnect all processes */
```





The MPI version of cpi

```
MPI_Init(&argc,&argv); /* Connect processes to each other */
MPI_Comm_size(MPI_COMM_WORLD,&numprocs); /* Get total number of processes */
MPI_Comm_rank(MPI_COMM_WORLD,&myid); /* Rank of this process */
MPI_Get_processor_name(name, &resultlen);

printf("This is Process-%d/%d running on %s \n",myid,numprocs,name);
MPI_Barrier(MPI_COMM_WORLD);


if(myid == 0) {
    printf("This program uses %d processes\n", numprocs);
    n = 1000000000;
}


MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); /* Send n from 0 to all others */

sum = 0.0;
h = 1.0/n;
for (i=myid+0.5; i<n; i+=numprocs) {
    sum += dx_arctan(i*h);
}
mypi = 4.0*h*sum;


/* Add all partial sums to each other and send to rank 0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if (myid == 0) {
    printf("pi is approximately %.16f\n",pi);
}
MPI_Finalize(); /* Disconnect all processes */
```






By Phonikal



By DaveSGage

Fast Fourier Transform



<http://2decomp.org>

Particle Data (N-Body)

