



University of Zurich^{UZH}

High Performance Computing Lecture 4

Douglas Potter

Cesare Cozza

Mark Eberlein

The Shell Prompt?

```
dhcp-94-215:~$ ssh ela
```

```
Last login: Mon May 30 14:16:07 2022 from dhcp-94-190.physik.uzh.ch
```

```
=====
IMPORTANT NOTICE FOR USERS of CScS facilities
```

```
Documentation: CScS User Portal - https://user.cscs.ch
```

```
Request support: https://support.cscs.ch
=====
```

```
[dpotter@ela1 ~]$ ssh eiger
```

```
Last login: Tue Oct 11 09:44:46 2022 from 148.187.1.10
```

```
=====
IMPORTANT NOTICE FOR USERS of Alps (EIGER)
```

```
Documentation: Alps User Guide - https://confluence.cscs.ch/x/\_gD0E
```

```
Request support: Service Desk at https://support.cscs.ch
=====
```

```
[eiger][dpotter@nid001080 ~]$
```

What is this?

“C”

Scripting

Regex

- C Crash Course (memory & pointers)
- Benefits of Scripting / Why Script?
 - Makes repetitive tasks faster
 - Reduces errors
 - Orchestration
- Focus on BASH scripting
 - Many other good tools exist (e.g., Python)
 - BASH works well for program Orchestration
 - BASH is omnipresent (on every Linux System)
 - Usually BASH is the default shell
 - You will encounter BASH in HPC
- Regular Expressions



C Data Types (signed or unsigned)

Type	Size in bits (bytes)	Description & Maximum Range (may be larger)
char	8 (1)	Smallest addressable unit on the machine [-127,+127] or [0,255]
short	16 (2)	At least [-32767,+32767] or [0,65535] (signed or unsigned)
int	16 (2) – 32 on daint	At least short, but may be long. Natural size of the machine
long	32 (4) – 64 on daint	[-2,147,483,647, +2,147,483,647] or [0, 4,294,967,295] (or larger)
long long	64 (8)	$[-2^{63} - 1, +2^{63} - 1]$ or $[0, 2^{64} - 1]$
float	32 (4)	IEEE 754 single-precision floating-point format
double	64 (8)	IEEE 754 double-precision floating-point format
long double	80, 96, 128 (10–16)	IEEE 754 quadruple-precision or other supported format



Fixed width integer types: `#include <stdint.h>`

Type	Size in bits (bytes)	Description & Maximum Range (may be larger)
int8_t	8 (1)	[-128,+127]
uint8_t	8 (1)	[0,255]
int16_t	16 (2)	[-32768,+32767]
uint16_t	16 (2)	[0,65535]
int32_t	32 (4)	[-2,147,483,648, +2,147,483,647]
uint32_t	32 (4)	[0, 4,294,967,295]
int64_t	64 (8)	$[-2^{63}, +2^{63} - 1]$
uint64_t	64 (8)	$[0, 2^{64} - 1]$



Better fixed width integer types?

Type	Size in bits (bytes)	Description & Maximum Range (may be larger)
int_least8_t	≥ 8 (1)	$[-128, +127]$
uint_least8_t	≥ 8 (1)	$[0, 255]$
int_least16_t	≥ 16 (2)	$[-32768, +32767]$
uint_least16_t	≥ 16 (2)	$[0, 65535]$
int_least32_t	≥ 32 (4)	$[-2,147,483,648, +2,147,483,647]$
uint_least32_t	≥ 32 (4)	$[0, 4,294,967,295]$
int_least64_t	≥ 64 (8)	$[-2^{63}, +2^{63} - 1]$
uint_least64_t	≥ 64 (8)	$[0, 2^{64} - 1]$



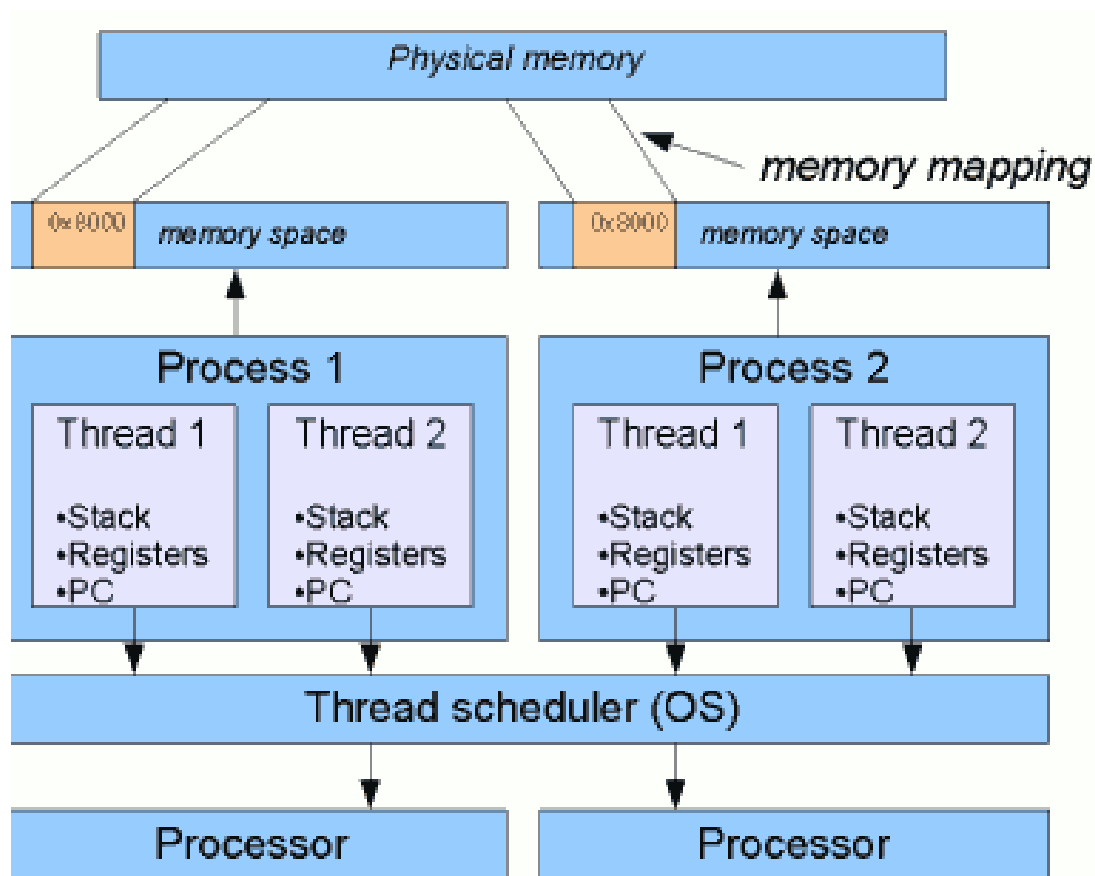
Even better fixed width integer types?

Type	Size in bits (bytes)	Description & Maximum Range (may be larger)
int_fast8_t	≥ 8 (1)	$[-128, +127]$
uint_fast8_t	≥ 8 (1)	$[0, 255]$
int_fast16_t	≥ 16 (2)	$[-32768, +32767]$
uint_fast16_t	≥ 16 (2)	$[0, 65535]$
int_fast32_t	≥ 32 (4)	$[-2,147,483,648, +2,147,483,647]$
uint_fast32_t	≥ 32 (4)	$[0, 4,294,967,295]$
int_fast64_t	≥ 64 (8)	$[-2^{63}, +2^{63} - 1]$
uint_fast64_t	≥ 64 (8)	$[0, 2^{64} - 1]$



Let's Talk Memory

- All processes have memory
- All variables are in memory
- Think of memory as an array
 - Starts at 0
 - Goes to $2^{64} - 1$
 - Not all indexes are valid (most not)
 - Each element is a "byte"
- Bigger types take more bytes and have consecutive addresses.





Basic Types

```
int i;
```

```
float f;
```

```
char c;
```

```
i = 12;
```

```
f = 23.88;
```

```
c = 'X';
```



Compound Type (structures)

```
struct person {  
    float height;  
    int age;  
};
```

```
struct person bob;  
bob.height = 1.91;  
bob.age = 24;
```



Pointers

C Code	Bytes	What is it?
<code>int i;</code>	4	An integer somewhere in memory. Compiler keeps track of where.
<code>int *p;</code>	8	A pointer to an integer. Contains the address of the integer.
<code>&i</code>	8	Pointer to (address of) the integer "i".
<code>p</code>	8	Address of an integer.
<code>*p</code>	4	What's at the address "p" (an integer)
<code>int **q</code>	8	A pointer to a pointer to an integer. e.g.: <code>q = &p</code>
<code>*q</code>	8	A pointer to an integer.
<code>**q</code>	4	An integer.



Structure Pointers

```
struct person {  
    float height;  
    int age;  
};
```

```
struct person bob;
```

```
bob.height = 1.91;  
bob.age = 24;
```

```
struct person grace;  
Struct person *p = &grace;
```

```
(*p).age = 40;
```

```
p->height = 1.79;
```



Dynamic Memory (allocation)

C Code	What it means
<code>float f;</code>	A single precision floating point number somewhere in memory
<code>float *p = &f;</code>	A variable that has the address of “f” (points to “f”)
<code>float *q;</code>	Another “pointer”. The value is undefined! Don’t use it (yet).
<code>q = malloc(4);</code>	Allocates four bytes and now “q” has the address. Don’t do this.
<code>q = malloc(sizeof(float));</code>	Same as before, but the compiler fills in the number of bytes needed.
<code>q = malloc(100*sizeof(float));</code>	Allocate 400 bytes (100 x 4) to store 100 float values.
<code>free(q);</code>	Free memory that “q” points to.



Pointer arithmetic

C Code	What it means
<code>float *p = malloc(100*sizeof(float));</code>	Allocate memory for 100 floating point numbers
<code>*p = 3.14;</code>	Set the “float” at p to 3.14. So set the first float of 100.
<code>p[0] = 8.1;</code>	Set the first float at p to 8.1.
<code>p[1] = -4.0;</code>	Set the second float to -4
<code>*(p+1) = -4.0;</code>	Same as above. Add “1 float” to the address “p” and dereference it.
<code>p[-1] = 666;</code>	Set the float before the memory block to 666. VERY BAD.
<code>float *q = p + 10;</code>	The pointer “q” now points to the 11th float in the memory.
<code>q[-2] = 555.5;</code>	The 9th float is set to 555.5. Valid, but uncommon.



The Golden Rule


```
float *p;
```

```
int i;
```

```
p[i] ≡ *(p+i)
```



Processes & Threads





top

```
top - 09:08:15 up 7 days, 22:56, 36 users, load average: 4.03, 2.63, 1.78
Tasks: 852 total, 4 running, 848 sleeping, 0 stopped, 0 zombie
%Cpu(s): 5.8 us, 2.3 sy, 0.4 ni, 90.4 id, 0.9 wa, 0.0 hi, 0.2 si, 0.0 st
KiB Mem: 26368408+total, 19931880+used, 64365276 free, 5413456 buffers
KiB Swap: 13421772+total, 5144 used, 13421257+free. 14725691+cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
30146	lyard	20	0	372268	248216	7768	R	147.06	0.094	4:23.17	python
4608	wchen	20	0	707304	307296	21864	R	94.118	0.117	2:42.02	pdos.py
22076	dcampi	20	0	139080	7484	4432	R	82.353	0.003	31:12.44	sshd
5549	dpotter	20	0	34592	4272	3068	R	5.882	0.002	0:00.01	top
25185	root	20	0	0	0	0	S	5.882	0.000	47:18.00	ptlrpcd_01_00
25186	root	20	0	0	0	0	S	5.882	0.000	43:56.62	ptlrpcd_01_01
25187	root	20	0	0	0	0	S	5.882	0.000	47:18.89	ptlrpcd_01_02
25188	root	20	0	0	0	0	S	5.882	0.000	43:50.86	ptlrpcd_01_03
25190	root	20	0	0	0	0	S	5.882	0.000	43:54.58	ptlrpcd_01_05
25193	root	20	0	0	0	0	S	5.882	0.000	47:13.45	ptlrpcd_01_08
25194	root	20	0	0	0	0	S	5.882	0.000	43:57.16	ptlrpcd_01_09
1	root	20	0	120628	6636	4008	S	0.000	0.003	35:12.94	systemd
2	root	20	0	0	0	0	S	0.000	0.000	0:31.76	kthreadd
3	root	20	0	0	0	0	S	0.000	0.000	0:37.09	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.000	0.000	0:00.00	kworker/0:0H
8	root	20	0	0	0	0	S	0.000	0.000	13:59.34	rcu_sched
9	root	20	0	0	0	0	S	0.000	0.000	0:00.00	rcu_bh
10	root	rt	0	0	0	0	S	0.000	0.000	0:05.07	migration/0
11	root	rt	0	0	0	0	S	0.000	0.000	0:01.84	watchdog/0



top -u munge

```
top - 09:09:45 up 7 days, 22:58, 36 users,  load average: 4.37, 3.08, 2.02
Tasks: 847 total,   4 running, 843 sleeping,   0 stopped,   0 zombie
%Cpu(s):  5.8 us,  2.3 sy,  0.4 ni, 90.4 id,  0.9 wa,  0.0 hi,  0.2 si,  0.0 st
KiB Mem: 26368408+total, 19914214+used, 64541940 free, 5413456 buffers
KiB Swap: 13421772+total,   5144 used, 13421257+free. 14730092+cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
18221	munge	20	0	2214036	5308	3404	S	0.000	0.002	1:16.34	munged



top -u munge -H

```
top - 09:10:29 up 7 days, 22:58, 36 users,  load average: 5.27, 3.45, 2.19
Threads: 1576 total,   7 running, 1569 sleeping,   0 stopped,   0 zombie
%Cpu(s): 12.8 us,  4.7 sy,  0.0 ni, 76.6 id,  5.6 wa,  0.0 hi,  0.3 si,  0.0 st
KiB Mem:  26368408+total, 19917673+used, 64507336 free,  5413456 buffers
KiB Swap: 13421772+total,   5144 used, 13421257+free. 14731705+cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
18221	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:18.11	munged
18222	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:02.86	munged
18223	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.70	munged
18224	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.66	munged
18225	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.75	munged
18226	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.72	munged
18227	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.71	munged
18228	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.74	munged
18229	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.70	munged
18230	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.72	munged
18231	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.71	munged
18232	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.74	munged
18233	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.76	munged
18234	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.71	munged
18235	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.70	munged
18236	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.73	munged
18237	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.74	munged
18238	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.72	munged
18239	munge	20	0	2214036	5308	3404	S	0.000	0.002	0:01.73	munged



Processes when you log in

```
dpotter@daint102:~> ps -ef f
```

UID	PID	PPID	C	STIME	TTY	STAT	TIME	CMD
root	1	0	0	Feb21	?	Ss	579:14	/sbin/init
root	40965	1	0	Mar26	?	Ss	0:22	/usr/sbin/sshd -D
root	138685	40965	0	Mar27	?	Ss	0:00	_ sshd: user [priv]
user	138688	138685	0	Mar27	?	S	0:00	_ sshd: user@notty
user	138689	138688	0	Mar27	?	Ss	0:00	_ /bin/someprog
root	27544	40965	0	18:31	?	Ss	0:00	_ sshd: dpotter [priv]
dpotter	27589	27544	0	18:31	?	S	0:00	_ sshd: dpotter@pts
dpotter	27590	27589	0	18:31	pts/21	Ss	0:00	_ -bash
dpotter	28822	27590	0	18:31	pts/21	R+	0:00	_ ps -ef f

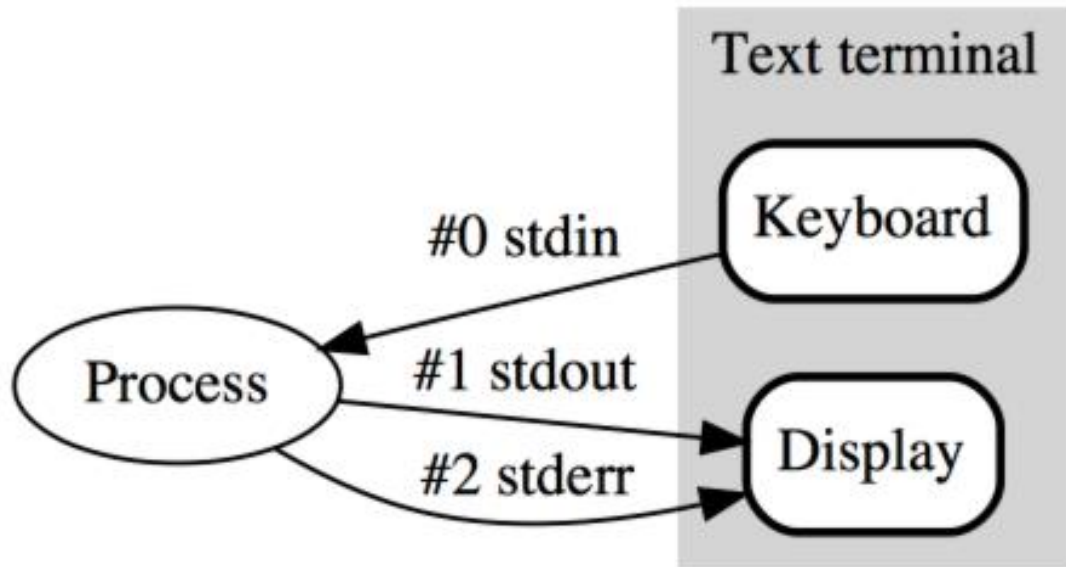
/dev/tty	“Teletype” Terminal
/dev/pts	Pseudo-Terminal (pty) slave



A Brief History of Shells

- RUNCOM: 1963 – run list of commands (legacy /etc/rc*)
- Multics Shell: 1965 – a command launcher
- Thompson Shell: Ken Thompson 1971 – “sh”
 - More of a “command interpreter”
 - Introduced redirection and pipes (more on this later)
- Bourne Shell: Stephen Bourne (Bell Labs) 1979
 - Designed as a “scripting language”
- BASH: Brian Fox (GNU Project) 1989 – “bash”
 - Combines features in the Bourne Shell, “csh” and “ksh”
 - POSIX (Portable Operating System Interface) compliant
 - Now the de facto standard
 - Stands for “Bourne Again SHell” (ha ha ha)

Standard Files





Example: “cat”

```
dpotter@daint102:~> cat
```

```
I typed this line.
```

```
I typed this line.
```

```
This is the second line.
```

```
This is the second line.
```

```
<Control-D>
```

```
dpotter@daint102:~>
```

```
dpotter@daint102:~> <Control-D>
```

```
dpotter@daint102:~> logout
```

```
Connection to daint102 closed.
```



Must be at the start of the line of input



Redirection



ela2:> **ls *.c**

blaio.c blas.c cpi.c foo.c hello2.c lapack.c mdl.c mpipi.c
res.c rz.c subarr.c walk2.c xthi.c



ela2:> **ls *.c >output.dat**
ela2:> **cat output.dat**



blaio.c
blas.c
cpi.c
foo.c
hello2.c
lapack.c
mdl.c
mpipi.c
res.c
rz.c
subarr.c
walk2.c
xthi.c



File Descriptors

File Descriptor	Normal Target	>output.dat
0 (stdin)	/dev/tty	/dev/tty
1 (stdout)	/dev/tty	output.dat
2 (stderr)	/dev/tty	/dev/tty
3 (first program file)	not used until a file is opened by the program	

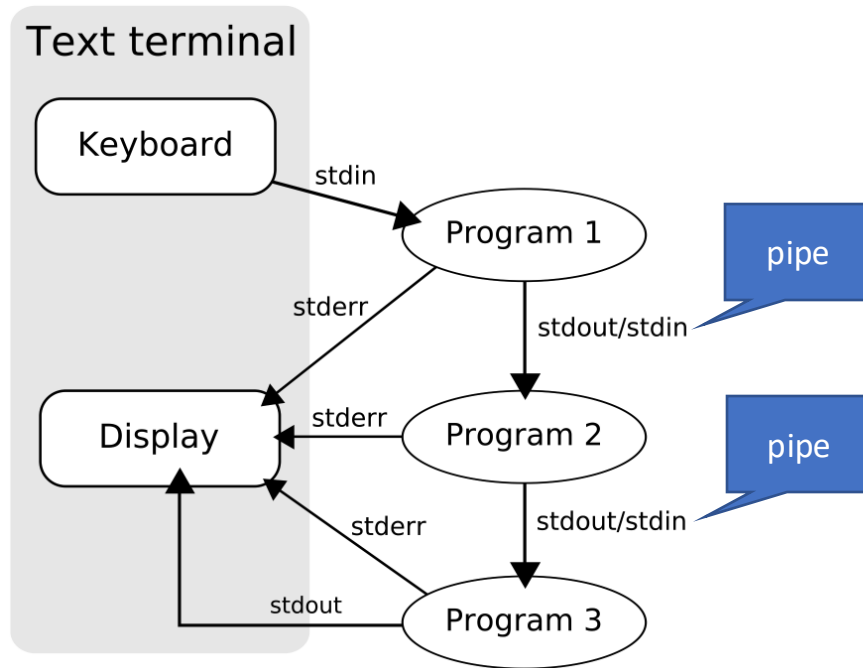
Redirect/To	stdin	stdout	stderr	"out.dat"
stdin	-	-	-	<out.dat
stdout	-	-	1>&2	>out.dat
stderr	-	2>&1	-	2>out.dat
stdout+stderr	-	-	-	&>out.dat

Appending to a file

```
ela3:> echo >data "This is one line"
ela3:> cat data
This is one line
ela3:> echo >data "This is line two"
ela3:> cat data
This is line two
ela3:> echo >>data "This is line three"
ela3:> cat data
This is line two
This is line three
```



Pipeline: program1 | program2 | program3





Example: “more”

```
dpotter@daint102:~/hpc1a> ps -ef | more
UID      PID    PPID  C  STIME TTY          TIME CMD
root         1        0  0  Feb21 ?        09:33:20 /sbin/init
root         2        0  0  Feb21 ?        00:02:24 [kthreadd]
root         3         2  0  Feb21 ?        00:11:35 [ksoftirqd/0]
root         5         2  0  Feb21 ?        00:00:00 [kworker/0:0H]
root         8         2  0  Feb21 ?        01:28:45 [rcu_sched]
root         9         2  0  Feb21 ?        00:00:00 [rcu_bh]
root        10         2  0  Feb21 ?        00:00:56 [migration/0]
root        11         2  0  Feb21 ?        00:00:16 [watchdog/0]
root        12         2  0  Feb21 ?        00:00:14 [watchdog/1]
root        13         2  0  Feb21 ?        00:01:00 [migration/1]
root        14         2  0  Feb21 ?        00:04:36 [ksoftirqd/1]
root        18         2  0  Feb21 ?        00:00:15 [watchdog/2]
root        19         2  0  Feb21 ?        00:00:54 [migration/2]
root        20         2  0  Feb21 ?        01:46:42 [ksoftirqd/2]
root        22         2  0  Feb21 ?        00:00:00 [kworker/2:0H]
root        23         2  0  Feb21 ?        00:00:14 [watchdog/3]
root        24         2  0  Feb21 ?        00:00:47 [migration/3]
root        25         2  0  Feb21 ?        00:04:13 [ksoftirqd/3]
root        27         2  0  Feb21 ?        00:00:00 [kworker/3:0H]
root        28         2  0  Feb21 ?        00:00:15 [watchdog/4]
root        29         2  0  Feb21 ?        00:01:01 [migration/4]
root        30         2  0  Feb21 ?        00:11:45 [ksoftirqd/4]
root        32         2  0  Feb21 ?        00:00:00 [kworker/4:0H]
--More--
```



Example: “less”

```
dpotter@daint102:~/hpc1a> ps -ef | less
UID      PID    PPID  C  STIME TTY          TIME CMD
root         1        0  0  Feb21 ?        09:33:20 /sbin/init
root         2        0  0  Feb21 ?        00:02:24 [kthreadd]
root         3         2  0  Feb21 ?        00:11:35 [ksoftirqd/0]
root         5         2  0  Feb21 ?        00:00:00 [kworker/0:0H]
root         8         2  0  Feb21 ?        01:28:45 [rcu_sched]
root         9         2  0  Feb21 ?        00:00:00 [rcu_bh]
root        10         2  0  Feb21 ?        00:00:56 [migration/0]
root        11         2  0  Feb21 ?        00:00:16 [watchdog/0]
root        12         2  0  Feb21 ?        00:00:14 [watchdog/1]
root        13         2  0  Feb21 ?        00:01:00 [migration/1]
root        14         2  0  Feb21 ?        00:04:36 [ksoftirqd/1]
root        18         2  0  Feb21 ?        00:00:15 [watchdog/2]
root        19         2  0  Feb21 ?        00:00:54 [migration/2]
root        20         2  0  Feb21 ?        01:46:42 [ksoftirqd/2]
root        22         2  0  Feb21 ?        00:00:00 [kworker/2:0H]
root        23         2  0  Feb21 ?        00:00:14 [watchdog/3]
root        24         2  0  Feb21 ?        00:00:47 [migration/3]
root        25         2  0  Feb21 ?        00:04:13 [ksoftirqd/3]
root        27         2  0  Feb21 ?        00:00:00 [kworker/3:0H]
root        28         2  0  Feb21 ?        00:00:15 [watchdog/4]
root        29         2  0  Feb21 ?        00:01:01 [migration/4]
root        30         2  0  Feb21 ?        00:11:45 [ksoftirqd/4]
root        32         2  0  Feb21 ?        00:00:00 [kworker/4:0H]
:
```

“less” is “more”

Who strive - you don't know how the others strive
To paint a little thing like that you smeared
Carelessly passing with your robes afloat,-
Yet do much less, so much less, Someone says,
(I know his name, no matter) - so much less!
Well, less is more, Lucrezia.

- *Andrea del Sarto*, 1855, Robert Browning



Example: My running processes

```
dpotter@daint102:~> ps -o pid,user,cmd -u dpotter f | less
```

```
PID USER      CMD
38794 dpotter  sshd: dpotter@pts/2
38795 dpotter  \_ -bash
41498 dpotter  \_ ps -o pid,user,cmd -u dpotter f
41499 dpotter  \_ less
... maybe other processes
```

(END)

```
daint104:> wc -l slurm-67584.out
144 slurm-67584.out
```

```
daint104:> less slurm-67584.out
daint104:> cat slurm-67584.out | less
          (obvious what happen here?)
```



Example: ls | wc

```
dpotter@daint102:~/hpc1a> ls
bad.c          cpi_openmp2.c  parse.pl
cpi            cpi_openmp.c   slurm-6188431.out
cpi.c         cpi_openmp.dat slurm-6188433.out
cpi.eps       cpi_openmp.job slurm-6192682.out
cpi_mpi       cpi.pdf        slurm-6312287.out
cpi_mpi.c     good.c         speedup.eps
cpi_mpi.job   good.s         speedup.gp
cpi_mpi.o     Makefile       speedup.pdf
cpi_openmp    old
```

```
dpotter@daint102:~/hpc1a> echo "just 3 words" | wc
1          3          13
```

```
dpotter@daint102:~/hpc1a> echo "X" | wc
1          1          2
```

```
dpotter@daint102:~/hpc1a> ls | wc
26         26        283
```



Mystery Solved?

```
dpotter@daint102:~/hpc1a> ls | cat
bad.c
cpi
cpi.c
cpi.eps
cpi_mpi
cpi_mpi.c
cpi_mpi.job
cpi_mpi.o
cpi_openmp
cpi_openmp2.c
cpi_openmp.c
cpi_openmp.dat
cpi_openmp.job
cpi.pdf
good.c
good.s
Makefile
old
parse.pl
slurm-6188431.out
slurm-6188433.out
slurm-6192682.out
slurm-6312287.out
speedup.eps
speedup.gp
speedup.pdf
```




Redirection

```
dpotter@daint102:~/hpc1a> ls >foo.dat
dpotter@daint102:~/hpc1a> cat foo.dat
bad.c
cpi
cpi.c
cpi.eps
cpi_mpi
cpi_mpi.c
cpi_mpi.job
cpi_mpi.o
cpi_openmp
cpi_openmp2.c
cpi_openmp.c
cpi_openmp.dat
cpi_openmp.job
cpi.pdf
foo.dat
good.c
good.s
Makefile
old
parse.pl
slurm-6188431.out
slurm-6188433.out
slurm-6192682.out
slurm-6312287.out
speedup.eps
speedup.gp
speedup.pdf
dpotter@daint102:~/hpc1a> wc <foo.dat
27 27 291
```



“Hello World” in C

```
#include <stdio.h>
```

```
int main(int argc, char *argv[]) {  
    printf("Hello world\n");  
    return 0;  
}
```



Arguments in C

```
#include <stdio.h>
```

```
int main(int argc, char *argv[]) {  
    int i;  
    for(i=0; i<argc; ++i)  
        printf("%d: %s\n", i, argv[i]);  
    return 0;  
}
```

Python

```
from sys import argv  
for (i,v) in enumerate(argv):  
    print("{}: {}".format(i,v))
```



Parameters

```
dpotter@daint103:~> ./hello2
0: ./hello2
dpotter@daint103:~> ./hello2 test a b c
0: ./hello2
1: test
2: a
3: b
4: c
dpotter@daint103:~> ./hello2 -h
0: ./hello2
1: -h
```



Wildcards (globbing)

```
daint103:~/hpc1a> ls *.c
bad.c  cpi.c  cpi_mpi.c  cpi_openmp2.c  cpi_openmp.c  good.c
daint103:~/hpc1a> ../hello2 *.c
0: ../hello2
1: bad.c
2: cpi.c
3: cpi_mpi.c
4: cpi_openmp2.c
5: cpi_openmp.c
6: good.c
daint103:~/hpc1a> ../hello2 "*.c"
0: ../hello2
1: *.c
daint103:~/hpc1a> ../hello2 '*.c'
0: ../hello2
1: *.c
daint103:~/hpc1a> ls "*.c"
ls: cannot access *.c: No such file or directory
```



Wildcards (globbing)

Wildcard	Description	Example
*	Matches zero or more	*.c (test.c example.c)
?	Matches exactly one	*.? (test.c foo.g)
[abc]	Only character in bracket	*.[ch]
[a-z]	Range of characters	[a-z]*, [a-zA-Z][0-9]*

man 7 glob

Expansion	Description	Example
{a,b,c}	Expands to the list of things	*.{c,h,cpp}
{s..e..i}	Counts from 's' to 'e' in increments of 'i'	{1..10..2}
~	Your home directory	~/bashrc
~+	Your current working directory	



Variables (simple)

```
hpc:~$ TEST1=12
hpc:~$ set | grep TEST1
TEST1=12
hpc:~$ TEST2=string
hpc:~$ set | grep TEST2
TEST2=string
hpc:~$ TEST3=a long string
-bash: long: command not found
hpc:~$ TEST3="a long string"
hpc:~$ set | grep TEST3
TEST3='a long string'
```



Variables (using them and spaces)

I am trying to construct a name in bash from a variable "I". If "I" is 12 I want the name to be "FOO12BAR", but if I try:
name="FOO\$IBAR" is does not work.

```
hpc:~$ echo $TEST1  
12
```

```
hpc:~$ echo $TEST3  
a long string
```

```
hpc:~$ echo $Test3
```

(nothing shows up here!)

```
hpc:~$ TEST4="a long string with spaces"
```

```
hpc:~$ echo ${TEST4}  
a long string with spaces
```

```
hpc:~$ echo this has space too  
this has space too
```

You're very close — the issue is that Bash doesn't know where the variable name (`I`) ends and the literal text (`BAR`) begins.

When you write:

```
bash
```

[Copy code](#)

```
name="FOO$IBAR"
```

Bash tries to interpret `$IBAR` as a single variable name, not `$I` followed by the string `BAR`.

To fix this, you need to use **braces** around the variable name so Bash knows where it ends:

```
bash
```

[Copy code](#)

```
name="FOO${I}BAR"
```

✓ Example:



Variables (quoting)

```
hpc:~$ echo "$TEST1"
```

```
12
```

```
hpc:~$ echo '$TEST1'
```

```
$TEST1
```

```
hpc:~$ echo "$TEST4"
```

```
a long string with spaces
```

```
hpc:~$ echo "Answer is $TEST1"
```

```
Answer is 12
```

```
daint103:~/hpc1a> ../hello2 string with space
```

```
0: ../hello2
```

```
1: string
```

```
2: with
```

```
3: space
```



Environment Variables

```
dpotter@daint102:~> ps -o pid,command -u dpotter f
```

```
PID COMMAND
```

```
87343 sshd: dpotter@pts/51
```

```
87345  \_ -bash
```

```
100054  \_ ps -o pid,command -u dpotter f
```

```
dpotter@daint102:~> bash
```

```
dpotter@daint102:~> ps -o pid,command -u dpotter f
```

```
PID COMMAND
```

```
87343 sshd: dpotter@pts/51
```

```
87345  \_ -bash
```

```
106380  \_ bash
```

```
107695  \_ ps -o pid,command -u dpotter f
```

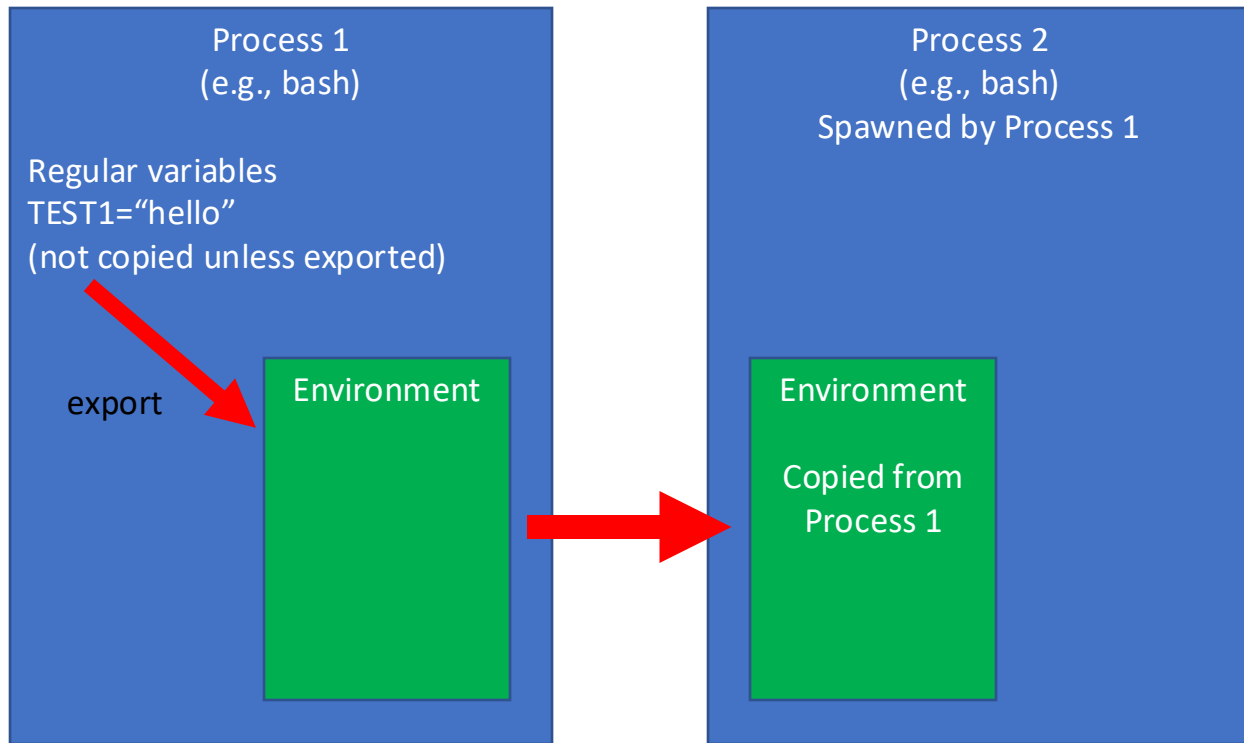


Environment Variables

```
dpotter@daint102:~> ETEST1="a variable"
dpotter@daint102:~> echo "$ETEST1"
a variable
dpotter@daint102:~> bash
dpotter@daint102:~> echo "$ETEST1"
(nothing shows up here!)
dpotter@daint102:~> exit
exit
dpotter@daint102:~> echo "$ETEST1"
a variable
dpotter@daint102:~> export ETEST1
dpotter@daint102:~> bash
dpotter@daint102:~> echo "$ETEST1"
a variable
dpotter@daint102:~> exit
```



Environment





Environment variable “PATH”

```
dpotter@daint102:~> printenv PATH
/apps/uzh/daint/bin:/apps/uzh/bin:/users/dpotter/.local/bin:/users/dpotter/local/bin:/apps/daint/UES/xalt
/0.7.6/bin:/opt/slurm/17.11.12.cscs/bin:/opt/cray/pe/mpt/7.7.2/gni/bin:/opt/cray/pe/perftools/7.0.2/bin:/
opt/cray/pe/papi/5.6.0.2/bin:/opt/cray/rca/2.2.18-6.0.7.0_33.3__g2aa4f39.ari/bin:/opt/cray/alps/6.6.43-
6.0.7.0_26.4__ga796da3.ari/sbin:/opt/cray/job/2.2.3-
6.0.7.0_44.1__g6c4e934.ari/bin:/opt/cray/pe/craype/2.5.15/bin:/opt/cray/pe/cce/8.7.3/binutils/x86_64/x86_
64-pc-linux-gnu/bin:/opt/cray/pe/cce/8.7.3/binutils/cross/x86_64-aarch64/aarch64-linux-
gnu/./bin:/opt/cray/pe/cce/8.7.3/utils/x86_64/bin:/opt/cray/pe/modules/3.2.10.6/bin:/opt/slurm/default/b
in:/apps/daint/system/bin:/apps/common/system/bin:/usr/local/bin:/usr/bin:/bin:/usr/bin/X11:/usr/lib/mit/
bin:/usr/lib/mit/sbin:/opt/cray/pe/bin:/apps/dora/system/sbin
dpotter@daint102:~> export PATH=""
dpotter@daint102:~> ls
-bash: ls: No such file or directory
dpotter@daint102:~> top
-bash: top: No such file or directory
dpotter@daint102:~> /usr/bin/ls
... normal output from ls appears here
```



Scripting (bash is a program)

```
dpotter@daint102:~> cat script.sh
echo "This is a script"
echo "This does not do very much"
date
dpotter@daint102:~> bash script.sh
This is a script
This does not do very much
Mon Apr 17 17:16:03 CEST 2023
dpotter@daint102:~> bash <script.sh
This is a script
This does not do very much
Mon Apr 17 17:16:18 CEST 2023
dpotter@daint102:~> cat script.sh | bash
This is a script
This does not do very much
Mon Apr 17 17:16:33 CEST 2023
```



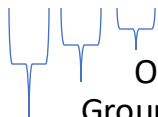
Scripting (Shebang)

```
dpotter@daint102:~> cat script.sh
#!/bin/bash
echo "This is a script"
echo "This does not do very much"
date
dpotter@daint102:~> chmod +x script.sh
dpotter@daint102:~> ./script.sh
This is a script
This does not do very much
Tue Apr 17 17:19:04 CEST 2018
dpotter@daint102:~> mv script.sh script
dpotter@daint102:~> ./script
This is a script
This does not do very much
Mon Apr 17 17:19:14 CEST 2023
```



File Permissions (a digression)

```
daint104:> ls -ld $HOME
drwx----- 59 dpotter uzh4 12288 Apr 18 09:08 /users/dpotter
daint104:> ls -l test_file
-rw-r--r-- 1 dpotter uzh0 0 Apr 18 10:14 test_file
daint104:> chmod o-r test_file
daint104:> ls -l test_file
-rw-r----- 1 dpotter uzh0 0 Apr 18 10:14 test_file
daint104:> chmod 755 test_file
daint104:> ls -l test_file
-rwxr-xr-x 1 dpotter uzh0 0 Apr 18 10:14 test_file
```


Other
Group
User (owner)

```
daint104:> chmod g-rx,o-rx test_file
daint104:> ls -l test_file
-rwx----- 1 dpotter uzh0 0 Apr 18 10:14 test_file
```




Shebang Generally

```
dpotter@daint102:~> cat output
```

```
#!/usr/bin/cat
```

A Shebang program will run the program listed on the first line, and it will "pipe" the file through it. Usually this is used for programs that take commands, but it is not restricted to this as you can see.

```
dpotter@daint102:~> ./output
```

```
#!/usr/bin/cat
```

A Shebang program will run the program listed on the first line, and it will "pipe" the file through it. Usually this is used for programs that take commands, but it is not restricted to this as you can see.



Parameters in bash scripts

```
dpotter@daint102:~> cat script
#!/bin/bash
echo "Parameter 1: $1"
echo "Parameter 2: $2"
dpotter@daint102:~> ./script
Parameter 1:
Parameter 2:
dpotter@daint102:~> ./script one two three
Parameter 1: one
Parameter 2: two
dpotter@daint102:~> ./script "one two three" four
Parameter 1: one two three
Parameter 2: four
dpotter@daint102:~> ls *.c
blaio.c blas.c cpi.c foo.c hello2.c lapack.c mdl.c mpipi.c
res.c rz.c subarr.c walk2.c xthi.c
dpotter@daint102:~> ./script *.c
Parameter 1: blaio.c
Parameter 2: blas.c
```



Parameters 0

```
dpotter@daint102:~> cat script0
#!/bin/bash
echo "Parameter 0: $0"
dpotter@daint102:~> ./script0
Parameter 0: ./script0
dpotter@daint102:~> ./script0 hello there
Parameter 0: ./script0
dpotter@daint102:~> cd code
dpotter@daint102:~/code> ../script0
Parameter 0: ../script0
dpotter@daint102:~> cd /tmp
dpotter@daint102:/tmp> /users/dpotter/script0
Parameter 0: /users/dpotter/script0

dpotter@daint102:~/code> script0
Parameter 0: /users/dpotter/.local/bin/script0
```



“simulate”

```
daint102:>./simulate  
How many stacks do you need?  
9  
What is the mean density?  
1.3  
Seed?  
314159365  
Grid size?  
256  
Generating stacks...  
done.  
Wallclock 1h45m12s
```



“simulate” Problems

- We would like to run this 100 times
 - Different set of parameters for each run
 - Independent jobs
- How can we feed in the variables?
 - Edit a file and pipe it in (or redirection)
`./simulate <parameters`
 - Create the file on-the-fly with “echo”
 - There is a better way though



“HERE” Documents

```
daint102:> cat script
#!/bin/bash
cat <<EOF
This is a line of text.
Here is another.
EOF
wc <<EOF
Input can go to any program.
We have three lines here.
And 14 words.
EOF
daint102:> ./script
This is a line of text.
Here is another.
3 14 69
```



“HERE” Documents

```
daint102:> cat script
#!/bin/bash
A="the variable A"
cat <<WHATEVER
Even more impressive is
that we can have variables.
```

```
Parameter 1: $1
Variable A: $A
WHATEVER
daint102:> ./script "parameter 1"
Even more impressive is
that we can have variables.
```

```
Parameter 1: parameter 1
Variable A: the variable A
```



“simulate” 1.0

```
#!/bin/bash
STACKS=9
DENSITY=1.3
SEED=314159265
GRID=256
./simulate <<EOD
$STACKS
$DENSITY
$SEED
$GRID
EOD
```




“simulate” 2.0

```
#!/bin/bash
STACKS=$1
DENSITY=$2
SEED=$3
GRID=$4
./simulate <<EOD
$STACKS
$DENSITY
$SEED
$GRID
EOD
```



The “if/then/else” statement


```
if <some condition> then  
    <take some action>  
end if
```

```
if (i < 100) {  
    printf("%d\n", i);  
}
```

```
if i < 100:  
    print(i)
```

BASH???

```
if $A < 100 then .
```





"If/then" in bash

```
if <command line>
then
    <command line(s)>
fi
```

```
if date
then
    echo "This was true???"
fi
dpotter@daint102:~> if date
> then
> echo "This was true???"
> fi
Tue Apr 17 17:34:25 CEST 2018
This was true???
```



“If/then” in bash

```
command1 ; command2 ; command3 ; ...
```

```
if <command line> ; then <command> ; fi
```

```
daint102:~> true
```

```
daint102:~> false
```

```
daint102:~> if true ; then echo yes ; fi  
yes
```

```
daint102:~> if false ; then echo yes ; fi
```

```
daint102:~>
```



“If/then/else if/else”

```
if <command line> ; then
    <command lines>
elif <command line> ; then
    <command lines>
elif <command line> ; then
    <command lines>
elif <command line> ; then
    <command lines>
else
    <command lines>
fi
```



“Hello World” in C

```
#include <stdio.h>
```

```
int main(int argc, char *argv[]) {  
    printf("Hello world\n");  
    return 0;  
}
```



The Return Code (error code)

```
daint102:~/hpc1a> ls >/dev/null
daint102:~/hpc1a> echo $?
0
daint102:~/hpc1a> ls >/dev/null no_such_file
ls: cannot access 'no_such_file': No such file or directory
daint102:~/hpc1a> echo $?
2
daint102:~/hpc1a> echo $?
0
daint102:~/hpc1a> true ; echo $?
0
daint102:~/hpc1a> false ; echo $?
1
```



“Boolean” Logic

```
daint102:> if true ; then echo yes ; fi
yes
daint102:> if true && true ; then echo yes ; fi
yes
daint102:> if true && false ; then echo yes ; fi
daint102:> if true || false ; then echo yes ; fi
yes
daint102:> if false || false ; then echo yes ; fi
daint102:> if true && false || true ; then echo yes ; fi
yes
daint102:> if true || false && true ; then echo yes ; fi
yes
daint102:> if true && echo second ; then echo yes ; fi
second
yes
daint102:> if true || echo second ; then echo yes ; fi
yes
```



Short circuit



test, [, [[

```
daint102:> A=12
daint102:> test $A -gt 10 ; echo $?
0
daint102:> test $A -gt 20 ; echo $?
1
daint102:> if test $A -gt 10 ; then echo large ; fi
large
daint102:> if [ $A -gt 10 ] ; then echo large ; fi
large
daint102:> if [[ $A -gt 10 ]] ; then echo large ; fi
large
daint102:> if [[ $A > 10 ]] ; then echo large ; fi
large
daint102:> if [[ $A > 20 ]] ; then echo large ; fi
daint102:> if [[ $A > 101 ]] ; then echo large ; fi
large
```

Lexicographical (Alphabetical)



“while”

```
while <command line>
do
    <command line>
done
daint102:> cat while.sh
#!/bin/bash
i=0
while test $i -lt 5 ; do
    echo $i
    let ++i
done
daint102:> ./while.sh
0
1
2
3
4
```



“for”

```
for NAME in word
do
    <command line>
done
daint104:> for V in a b c ; do echo $V ; done
a
b
c
daint104:> ls slurm-6*
slurm-67584.out  slurm-67594.out  slurm-67640.out  slurm-
6878386.out
daint104:> for NAME in slurm-6* ; do echo $NAME ; done
slurm-67584.out
slurm-67594.out
slurm-67640.out
slurm-6878386.out
daint104:> for (( i=0; i<10 ; ++i )) ; do echo $i ; done
0
1
...
9
```



“shift”

```
daint102:> cat shift
#!/bin/bash
while test -n "$1" ; do
    echo "Processing $1"
    shift
done
daint102:> ./shift
daint102:> ls *.c
blaio.c blas.c cpi.c foo.c hello2.c lapack.c mdl.c
mpipi.c res.c rz.c subarr.c walk2.c xthi.c
daint102:> ./shift *.c
Processing blaio.c
Processing blas.c
Processing cpi.c
Processing foo.c
Processing hello2.c
Processing lapack.c
Processing mdl.c
Processing mpipi.c
Processing res.c
Processing rz.c
Processing subarr.c
Processing walk2.c
Processing xthi.c
```



Integer Arithmetic $\$(())$ – also $(())$

```
ela3:~> A=90
ela3:~> B=$((A/2+7))
ela3:~> echo $B
52
ela3:~> C=$((B*2+A))
ela3:~> echo $C
194
ela3:~> D=$((C/5))
ela3:~> echo $D
38
```



Real Arithmetic ???

```
ela3:~> A=$((30.2))  
-bash: 30.2: syntax error: invalid arithmetic operator (error  
token is ".2")  
ela3:~> A=30.2  
ela3:~> B=12.4  
ela3:~> C=$((A+B))  
-bash: 30.2: syntax error: invalid arithmetic operator (error  
token is ".2")
```

Well, drat!



Real Arithmetic – not possible?

```
daint103:> bc
bc 1.06.95
Copyright 1991-1994, 1997, 1998, 2000, 2004, 2006 Free Software
Foundation, Inc.
This is free software with ABSOLUTELY NO WARRANTY.
For details type `warranty'.
1/3
0
scale=10
1/3
.3333333333
scale=9;1/3
.333333333
daint103:> echo "scale=10;1.0/3.0" | bc
.3333333333
```



Output Capture \$(), ``

```
daint103:> ANSWER=$(echo "scale=10;1.0/3.0" | bc)
daint103:> echo $ANSWER
.3333333333
daint103:> RESULT=`echo "scale=10;1.0/3.0" | bc`
daint103:> echo $RESULT
.3333333333
daint103:> ls *.c
blaio.c blas.c cpi.c foo.c hello2.c lapack.c
mdl.c mpipi.c res.c rz.c subarr.c walk2.c xthi.c
daint103:> FILES=$(ls *.c)
daint103:> echo $FILES
blaio.c blas.c cpi.c foo.c hello2.c lapack.c mdl.c mpipi.c
res.c rz.c subarr.c walk2.c xthi.c
```




Functions

```
daint103:> cat distance
#!/bin/bash
function distance () {
    local x=$1
    local y=$2
    echo "scale=10;sqrt($x*$x + $y*$y)" | bc
    return 0
}
distance 3 4
distance 10 9
daint103:> ./distance
5.0000000000
13.4536240470
```



Function Capture

```
daint103:> cat distance
#!/bin/bash
function distance () {
    local x=$1
    local y=$2
    echo "scale=10;sqrt($x*$x + $y*$y)" | bc
}
D1=$(distance 3 4)
D2=$(distance 10 9)
echo "Distance 1 is $D1"
daint103:> ./distance
Distance 1 is 5.0000000000
```



“simulate” revisited

```
daint103:> cat script.job
#!/bin/bash
function distance () {
    local x=$1
    local y=$2
    echo "scale=10;sqrt($x*$x + $y*$y)" | bc
}
cat <<EOF
9
$(distance 3 4)
314159265
$(distance 10 9)
EOF
ela3:> ./script.job
9
5.0000000000
314159265
13.4536240470
```

“cat”, but could
be “simulate”



Functions return values

```
daint103:> cat prime
#!/bin/bash
function isprime () {
    # should check if $1 is prime, but:
    return 0
}
if isprime $1 ; then
    echo "$1 is prime"
else
    echo "$1 is not prime"
fi
daint103:> ./prime 12
12 is prime
```

Advanced Features

- BASH has arrays as well
- String Manipulation

[←](#) [→](#) [↺](#) [Secure](#) | <https://www.tldp.org/LDP/abs/html/string-manipulation.html>

[Prev](#)

Advanced Bash-Scripting Guide:
Chapter 10. Manipulating Variables

10.1. Manipulating Strings

Bash supports a surprising number of string manipulation operations. Unfortunately, these tools lack a unified syntax, and others fall under the functionality of the UNIX [expr](#) command. This results in inconsistent command syntax across different shells.

String Length

`${#string}`

`expr length $string`

These are the equivalent of `strlen()` in C.

`expr "$string" : '.*'`

```
stringZ=abcABC123ABCabc
echo ${#stringZ}           # 15
echo `expr length $stringZ` # 15
echo `expr "$stringZ" : '.*'\` # 15
```

Example 10-1. Inserting a blank line between paragraphs in a text file



Grep: Searching in files

```
hpc:~$ cat interim.txt  
Interim  
Lola Ridge, 1873
```

```
The earth is motionless  
And poised in space ...  
A great bird resting in its flight  
Between the alleys of the stars.  
It is the wind's hour off ...  
The wind has nestled down among the corn ...  
The two speak privately together,  
Awaiting the whirr of wings.  
hpc:~$ grep alley interim.txt  
Between the alleys of the stars.
```



Grep: Searching in files

```
hpc:~$ grep --color alley interim.txt
Between the alleys of the stars.
hpc:~$ grep --color the interim.txt
Between the alleys of the stars.
It is the wind's hour off ...
The wind has nestled down among the corn ...
The two speak privately together,
Awaiting the whirr of wings.
hpc:~$ grep -c the interim.txt
5
hpc:~$ grep -q the interim.txt
hpc:~$ echo $?
0
```



Grep: Searching in files

```
hpc:~$ if grep -q alley interim.txt ; then
> echo found
> fi
found
hpc:~$ if grep -q apple interim.txt ; then
> echo found
> fi
hpc:~$ grep -q apple interim.txt
hpc:~$ echo $?
1
```




Grep: Some useful flags

```
hpc:~$ grep --color and interim.txt
hpc:~$ grep --color -i and interim.txt
And poised in space ...
hpc:~$ grep --color -i -A 1 and interim.txt
And poised in space ...
A great bird resting in its flight
hpc:~$ grep --color -i -B 1 and interim.txt
The earth is motionless
And poised in space ...
hpc:~$ grep --color -i -C 1 and interim.txt
The earth is motionless
And poised in space ...
A great bird resting in its flight
```



Global Regular Expression Print

- Appear in Version 4 Unix in 1973
- Created by Ken Thompson (original Unix)
- Based on a regular expression parser in “ed”
- “g/re/p” command
 - /re/ – Search for a regular expression “re”
 - g – do it globally
 - p – print the result
- Hence: g/re/p → grep!



Global Regular Expression Print

```
hpc:~$ grep --color a..ey interim.txt
Between the alleys of the stars.
hpc:~$ grep --color all. interim.txt
Between the alleys of the stars.
hpc:~$ grep --color all.* interim.txt
Between the alleys of the stars.
hpc:~$ grep --color al*ey interim.txt
Between the alleys of the stars.
hpc:~$ grep --color we*n interim.txt
Between the alleys of the stars.
The wind has nestled down among the corn ...
```



Extended G R E P

```
hpc:~$ egrep --color we*n interim.txt
Between the alleys of the stars.
The wind has nestled down among the corn ...
hpc:~$ egrep --color we+n interim.txt
Between the alleys of the stars.
hpc:~$ egrep --color wee*n interim.txt
Between the alleys of the stars.
hpc:~$ egrep --color we?x?e?v?n interim.txt
Between the alleys of the stars.
The wind has nestled down among the corn ...
```



Extended G R E P

```
hpc:~$ export GREP_OPTIONS=--color
```

```
hpc:~$ egrep -i the interim.txt
```

The earth is motionless

Between the alleys of the stars.

It is the wind's hour off ...

The wind has nestled down among the corn ...

The two speak privately together,

Awaiting the whirr of wings.

```
hpc:~$ egrep -i 'the wind' interim.txt
```

It is the wind's hour off ...

The wind has nestled down among the corn ...



Greedy Matching (the default)

```
hpc:~$ egrep 'alleys' interim.txt
```

Between the **alleys** of the stars.

```
hpc:~$ egrep 'all.*s' interim.txt
```

Between the **alleys of the stars**.

Greedy Match

```
hpc:~$ egrep 'all.*?s' interim.txt
```

Between the **alleys** of the stars.

Minimum Match



Alternate values (or)

```
hpc:~$ export GREP_OPTIONS='-i --color'
hpc:~$ egrep 'the wind' interim.txt
It is the wind's hour off ...
The wind has nestled down among the corn ...
hpc:~$ egrep 'the earth' interim.txt
The earth is motionless
hpc:~$ egrep 'the wind|the earth' interim.txt
The earth is motionless
It is the wind's hour off ...
The wind has nestled down among the corn ...
```



Parenthesis

```
hpc:~$ export GREP_OPTIONS='-i --color'
hpc:~$ egrep 'the wind|earth' interim.txt
The earth is motionless
It is the wind's hour off ...
The wind has nestled down among the corn ...
hpc:~$ egrep 'the (wind|earth)' interim.txt
The earth is motionless
It is the wind's hour off ...
The wind has nestled down among the corn ...
```




regex101.com

The screenshot shows the regex101.com website. The 'REGULAR EXPRESSION' field contains the pattern `/ the`. The 'TEST STRING' field contains a poem snippet: 'Interim Lola Ridge, 1873 The earth is motionless And poised in space ... A great bird resting in its flight Between the alleys of the stars. It is the wind's hour off ... The wind has nestled down among the corn ... The two speak privately together, Awaiting the whirr of wings.' The 'EXPLANATION' panel on the right shows 9 matches for the pattern 'the' in the test string, with details like 'Full match' and character ranges (e.g., 26-29 for 'The'). The 'QUICK REFERENCE' panel at the bottom right lists various regex tokens and their symbols.

REGULAR EXPRESSION	TEST STRING	EXPLANATION
<code>/ the</code>	Interim Lola Ridge, 1873 The earth is motionless And poised in space ... A great bird resting in its flight Between the alleys of the stars. It is the wind's hour off ... The wind has nestled down among the corn ... The two speak privately together, Awaiting the whirr of wings.	MATCH INFORMATION Match 1 Full match 26-29 'The' Match 2 Full match 117-120 'the' Match 3 Full match 131-134 'the' Match 4 Full match 148-151 'the' Match 5 Full match 172-175 'The' Match 6 Full match 204-207 'the' Match 7 Full match 217-220 'The' Match 8 Full match 245-248 'the' Match 9 Full match 260-263 'the'

QUICK REFERENCE
A single character of: a, b or c
A character except: a, b or c
A character in the range: a-z
A character not in the range: a-z
A character in the range: a-z or A-Z
Any single character
Any whitespace character
Any non-whitespace character
Any digit
Any non-digit
Any word character
Any non-word character
Capture everything enclosed
Match either a or b
Zero or one of a
Zero or more of a
One or more of a
Exactly 3 of a
3 or more of a
Between 3 and 6 of a
Start of string
End of string
A word boundary
Non-word boundary

QUICK REFERENCE		
Search reference		
all tokens	A single character of: a, b or c	[abc]
common tokens	A character except: a, b or c	[^abc]
general tokens	A character in the range: a-z	[a-z]
anchors	A character not in the range: a-z	[^a-z]
meta sequences	A character in the range: a-z or A-Z	[a-zA-Z]
quantifiers	Any single character	.
group constructs	Any whitespace character	\s
character classes	Any non-whitespace character	\S
flags/modifiers	Any digit	\d
substitution	Any non-digit	\D
	Any word character	\w
	Any non-word character	\W
	Capture everything enclosed	(...)
	Match either a or b	(a b)
	Zero or one of a	a?
	Zero or more of a	a*
	One or more of a	a+
	Exactly 3 of a	a{3}
	3 or more of a	a{3,}
	Between 3 and 6 of a	a{3,6}
	Start of string	^
	End of string	\$
	A word boundary	\b
	Non-word boundary	\B



BASH Regular Expressions

```
hpc:~$ cat retest.sh
#!/bin/bash
while read A ; do
    echo "$A"
done
hpc:~$ ./retest.sh <interim.txt
Interim
Lola Ridge, 1873
```

```
The earth is motionless
And poised in space ...
A great bird resting in its flight
Between the alleys of the stars.
It is the wind's hour off ...
The wind has nestled down among the corn ...
The two speak privately together,
Awaiting the whirr of wings.
```



BASH Regular Expressions

```
hpc:~$ cat retest.sh
#!/bin/bash
while read A ; do
    if [[ "$A" =~ we*n ]] ; then
        echo "$A"
    fi
done
hpc:~$ ./retest.sh <interim.txt
Between the alleys of the stars.
The wind has nestled down among the corn ...
hpc:~$ [[ "Hello world" =~ music ]] ; echo $?
1
hpc:~$ [[ "Hello world" =~ ll ]] ; echo $?
0
```



BASH Regular Expressions

```
hpc:~$ cat report.txt
Starting calculations
Phase 1 complete, 3.11 seconds
Phase 2 complete, 10.92 seconds
Phase 3 complete, 1.15 seconds
Calculations complete
hpc:~$ egrep 'Phase [0-9] complete, [0-9]+\.[0-9]+ seconds' report.txt
Phase 1 complete, 3.11 seconds
Phase 2 complete, 10.92 seconds
Phase 3 complete, 1.15 seconds
```



BASH Regular Expressions

```
hpc:~$ cat retest.sh
#!/bin/bash
RE="Phase [0-9] complete, [0-9]+\.[0-9]+
seconds"
while read A ; do
    if [[ "$A" =~ $RE ]] ; then
        echo "$A"
    fi
done
hpc:~$ ./retest.sh <report.txt
Phase 1 complete, 3.11 seconds
Phase 2 complete, 10.92 seconds
Phase 3 complete, 1.15 seconds
```



Regular Expression Capture

```
hpc:~$ cat retest.sh
#!/bin/bash
RE="Phase ([0-9]) complete, ([0-9]+\.[0-9]+)
seconds"
while read A ; do
    if [[ "$A" =~ $RE ]] ; then
        echo "Match: ${BASH_REMATCH[0]}"
        echo "Time:  ${BASH_REMATCH[2]}"
    fi
done
hpc:~$ ./retest.sh <report.txt
Match: Phase 1 complete, 3.11 seconds
Time: 3.11
Match: Phase 2 complete, 10.92 seconds
Time: 10.92
Match: Phase 3 complete, 1.15 seconds
Time: 1.15
```



Group matches

Online regex tester and debug: x

Secure | https://regex101.com

regular expressions 101 @regex101 donate contact bug reports & feedback wiki

REGULAR EXPRESSION 3 matches, 102 steps (~0ms)

/ Phase ([0-9]) complete, ([0-9]+\.[0-9]+) seconds /g

TEST STRING SWITCH TO UNIT TESTS

Starting calculations
Phase 1 complete, 3.11 seconds
Phase 2 complete, 10.92 seconds
Phase 3 complete, 1.15 seconds
Calculations complete

EXPLANATION

MATCH INFORMATION

Match 1

Full match 22-52 `Phase 1 complete, 3.11 seconds`

Group 1. 28-29 `1`

Group 2. 40-44 `3.11`

Match 2

Full match 53-84 `Phase 2 complete, 10.92 seconds`

Group 1. 59-60 `2`

Group 2. 71-76 `10.92`

Match 3

Full match 85-115 `Phase 3 complete, 1.15 seconds`

Group 1. 91-92 `3`

Group 2. 103-107 `1.15`

SUBSTITUTION

EXPLANATION

MATCH INFORMATION

QUICK REFERENCE

Search reference

- All Tokens
- ★ Common Tokens ✓
- General Tokens
- ⚓ Anchors
- ⚙ Meta Sequences
- * Quantifiers
- ⦿ Group Constructs
- 📦 Character Classes
- 🚩 Flags/Modifiers
- 🔗 Substitution

A single character... [abc]

A character exce... [^abc]

A character in the... [a-z]

A character not i... [^a-z]

A character in... [a-zA-Z]

Any single character .

Any whitespace chara... \s

Any non-whitespace c... \S

Any digit \d

Any non-digit \D

Any word character \w

Any non-word charac... \W

Capture everythin... (...)

Match either a or b (a|b)

Zero or one of a a?

Zero or more of a a*

One or more of a a+

Exactly 3 of a a{3}

3 or more of a a{3,}

Between 3 and 6... a{3,6}

Start of string ^

End of string \$

A word boundary \b

Non-word boundary \B



Regular Expression Capture

```
hpc:~$ cat retest.sh
#!/bin/bash
RE="Phase ([0-9]) complete, ([0-9]+\.[0-9]+)
seconds"
TIME=0
while read A ; do
    if [[ "$A" =~ $RE ]] ; then
        TIME=$((echo "scale=5;$TIME +\
${BASH_REMATCH[2]}" | bc))
    fi
done
echo "Total time: $TIME seconds"
hpc:~$ ./retest.sh <report.txt
Total time: 15.18 seconds
```




Python

```
hpc:~$ cat retest.py
#!/usr/bin/python
import re
import fileinput
pattern = re.compile(
    'Phase ([0-9]) complete, ([0-9]+\.[0-9]+) seconds')
time = 0
for line in fileinput.input():
    match = pattern.match(line) # None or "Match Object"
    if match:
        time += float(match.group(2))
print("Total time: %f seconds" % time)
hpc:~$ ./retest.py <report.txt
Total time: 15.180000 seconds
```